


Module 6 :

Reinforcement Learning

國立中山大學
資訊工程系
張雲南





5

Reinforcement Learning



Application

◆ Autonomous helicopter

- ◆ <http://heli.stanford.edu/>

- ◆ <https://www.youtube.com/watch?v=VCdxqn0fcnE>

◆ Robotics

- ◆ Towards Learning Robot Table Tennis

- ◆ <https://www.youtube.com/watch?v=SH3bADiB7uQ>

- ◆ Guided Cost Learning: Deep Inverse Optimal Control via Policy Optimization

- ◆ <https://www.youtube.com/watch?v=hXxaepw0zAw>

◆ Save power

- ◆ Transforming Cooling Optimization for Green Data Center via Deep Reinforcement Learning

◆ Game

- ◆ AlphaGo


Application

◆ NLP

- ◆ IBM Watson



◆ Stock trading

- ◆ https://github.com/llSourceCell/Reinforcement_Learning_for_Stock_Prediction



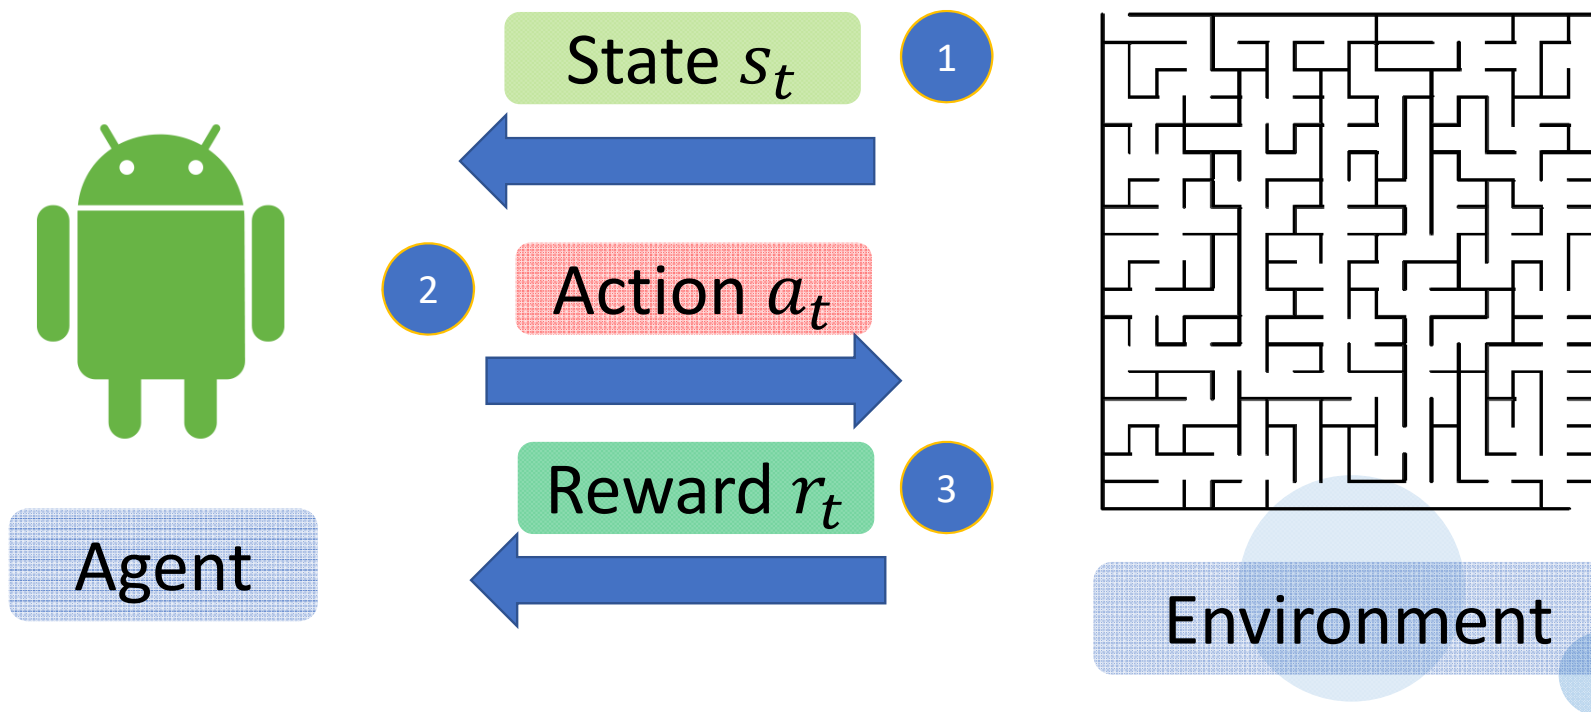
OpenAI gym

- ◆ A toolkit for developing and comparing reinforcement learning algorithms.




Reinforcement Learning


- ◆ Train an **Agent/Actor** to take actions that can receive best **total reward**.
 - It involves a whole sequence of actions.



Total Reward - Tetris




Total Reward - Tetris



Next Block


If you choose to go straight down, one line can be cleared.

Total Reward - Tetris




If you go straight down, you can get a reward of **1 line** for your “straight-down” action .

Total Reward - Tetris



But if you **rotate** it, you don't get reward in this action.


Total Reward - Tetris



Next Block


If you rotate it, you don't get reward for this action.

Total Reward - Tetris




But you can get a higher reward (**3 lines**) in the next action.


Total Reward - Tetris




Reward : 1



Reward : 0



Reward : 3



Reinforcement Learning


◆ In Gomoku, one who can form an unbroken chain of five is the winner

State s_t

Board

Environment

Opponent



Action a_t

Where to place the stone


Reward r_t

Win : get 1 point

Not-Win : get 0

Loss: get -1 point

Reward



Autonomous helicopter



Agent

Wind

Gravity

Motor control

-10/0/10/100

State

Action

Reward



Environment

Reward Grading
System

Autonomous helicopter

State

Action

Reward

Episode 1

◆ Our goal is to fly a circle (blue line).



Autonomous helicopter

State

Action

Reward


Episode 4

◆ Our goal is to fly a circle (blue line).



Mountain Car Problem

- ◆ Mountain Car is a problem in which an under-powered car must drive up a steep hill.
- ◆ Since gravity is stronger than the car's engine, even at full throttle, the car cannot simply accelerate to climb up the steep slope.



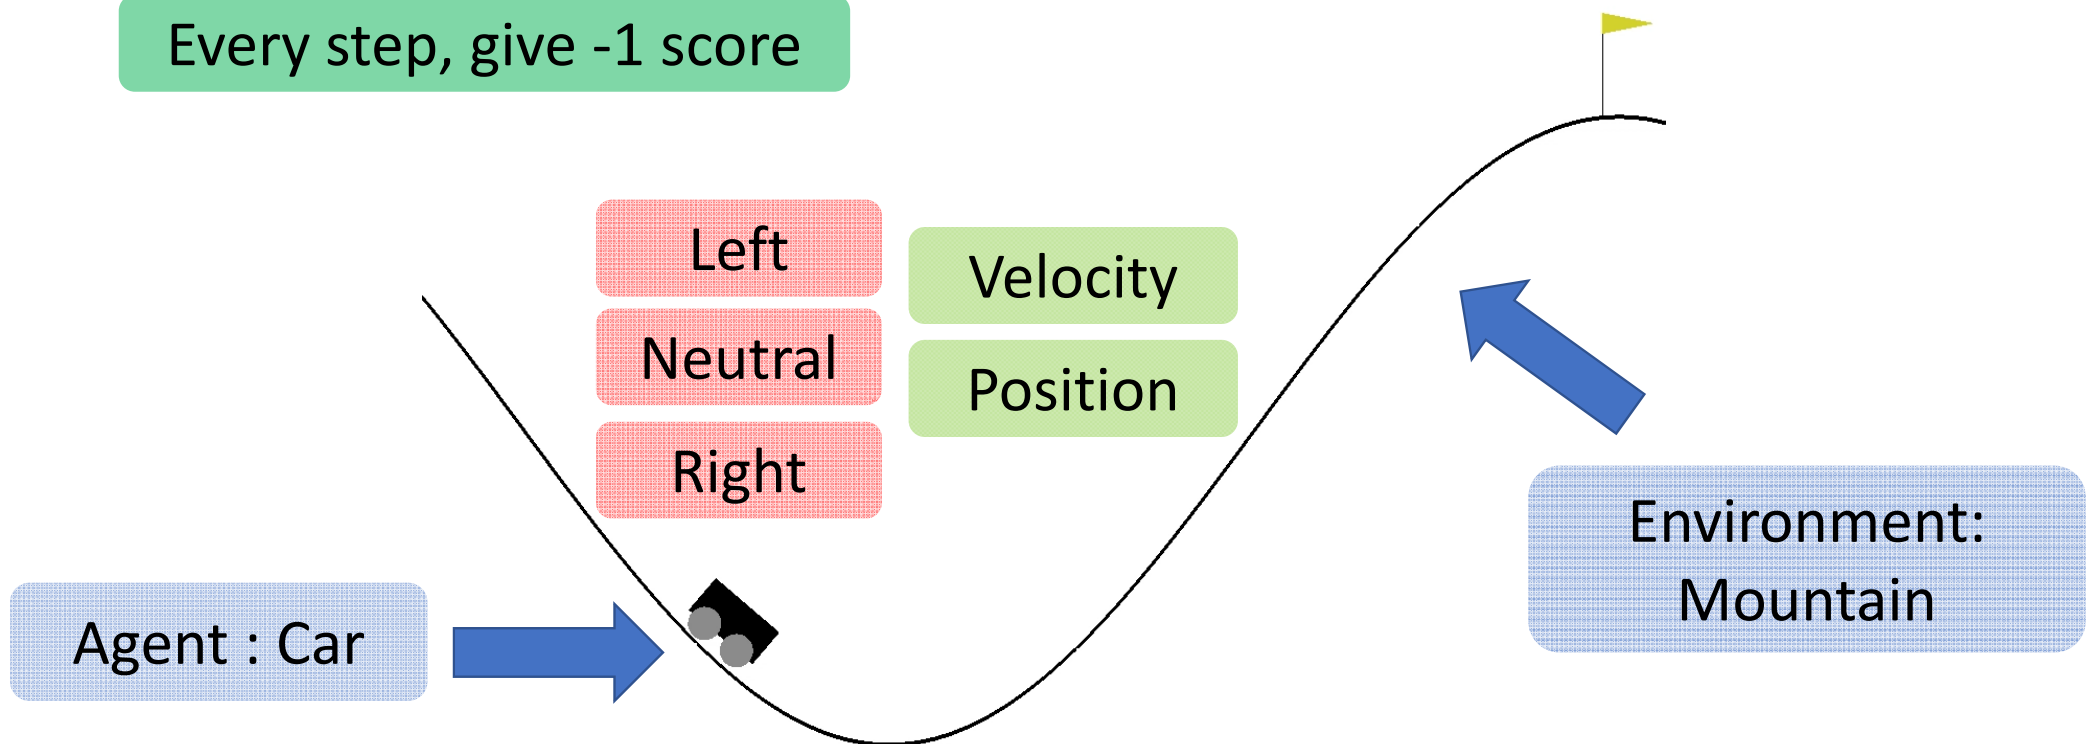
Mountain Car Problem

State

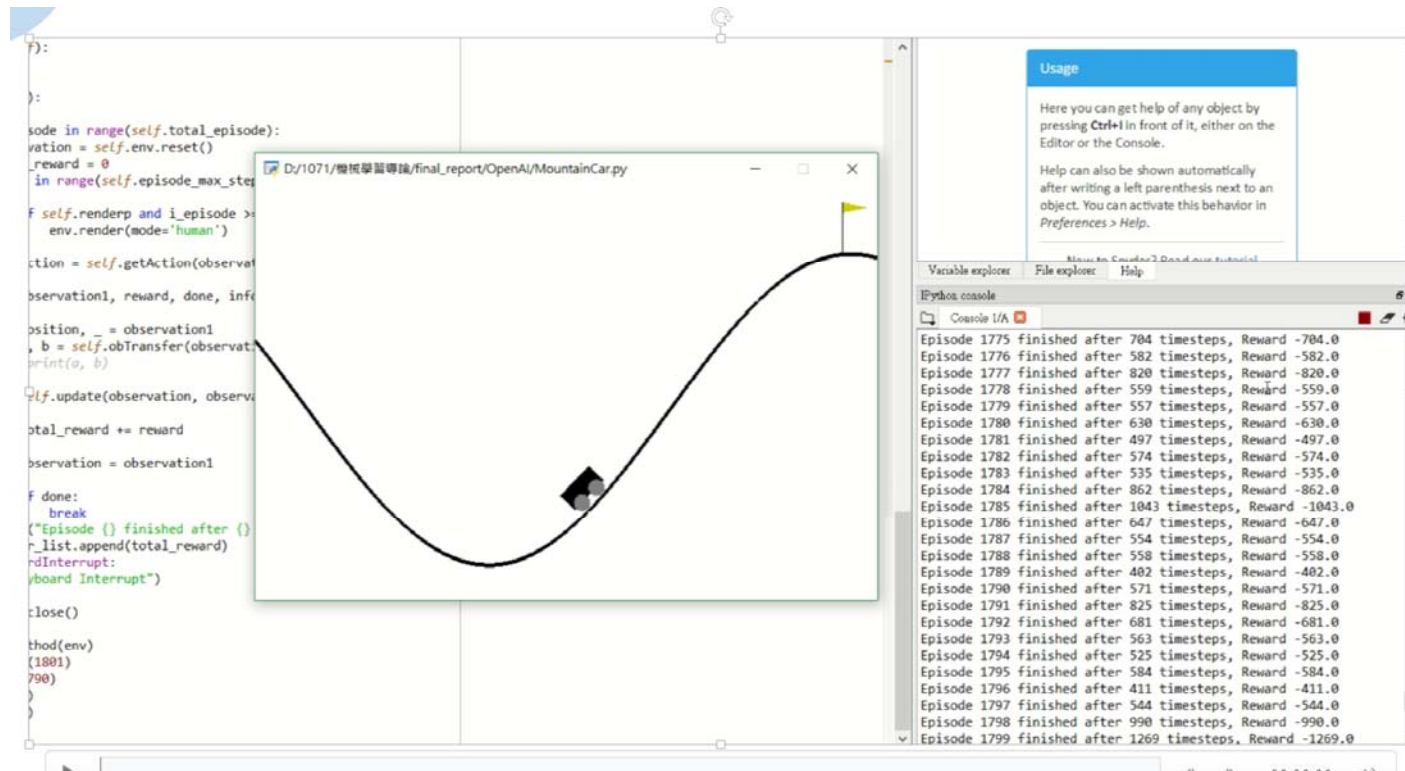
Action

Reward

Every step, give -1 score



Mountain Car in OpenAI gym



Text generator


State

Action

Reward

Why did he
talk that?

How old are you?



Environment:
person


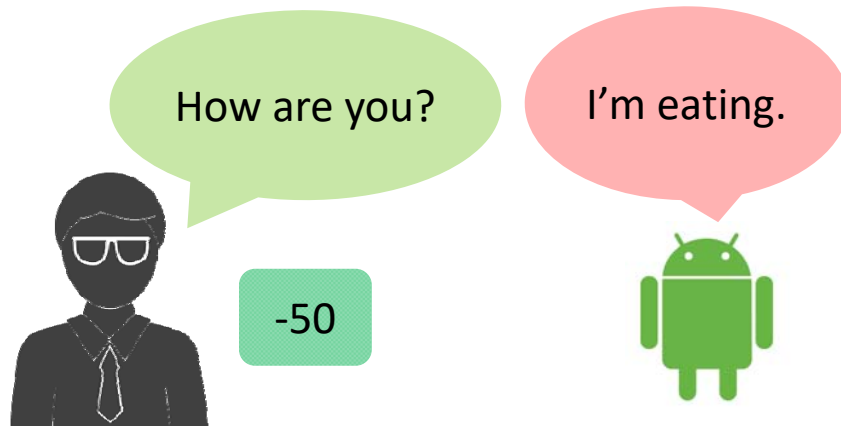


I'm eating.



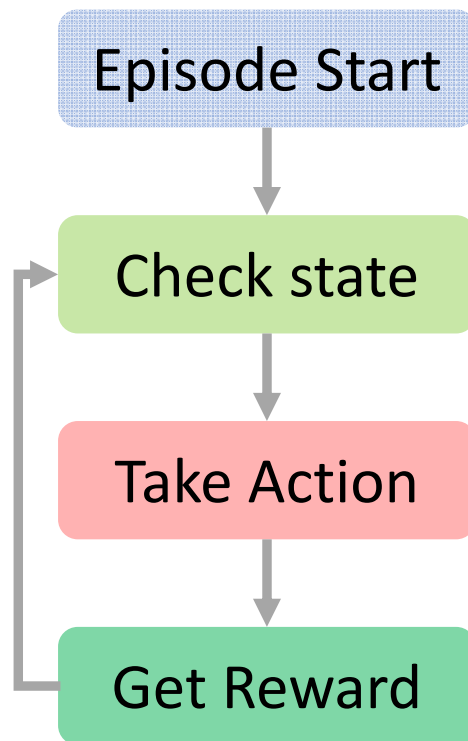
Agent

Next time, I
should not say
this.


Text generator



Learning Process




Episode Start




Get Reward

1

Check state



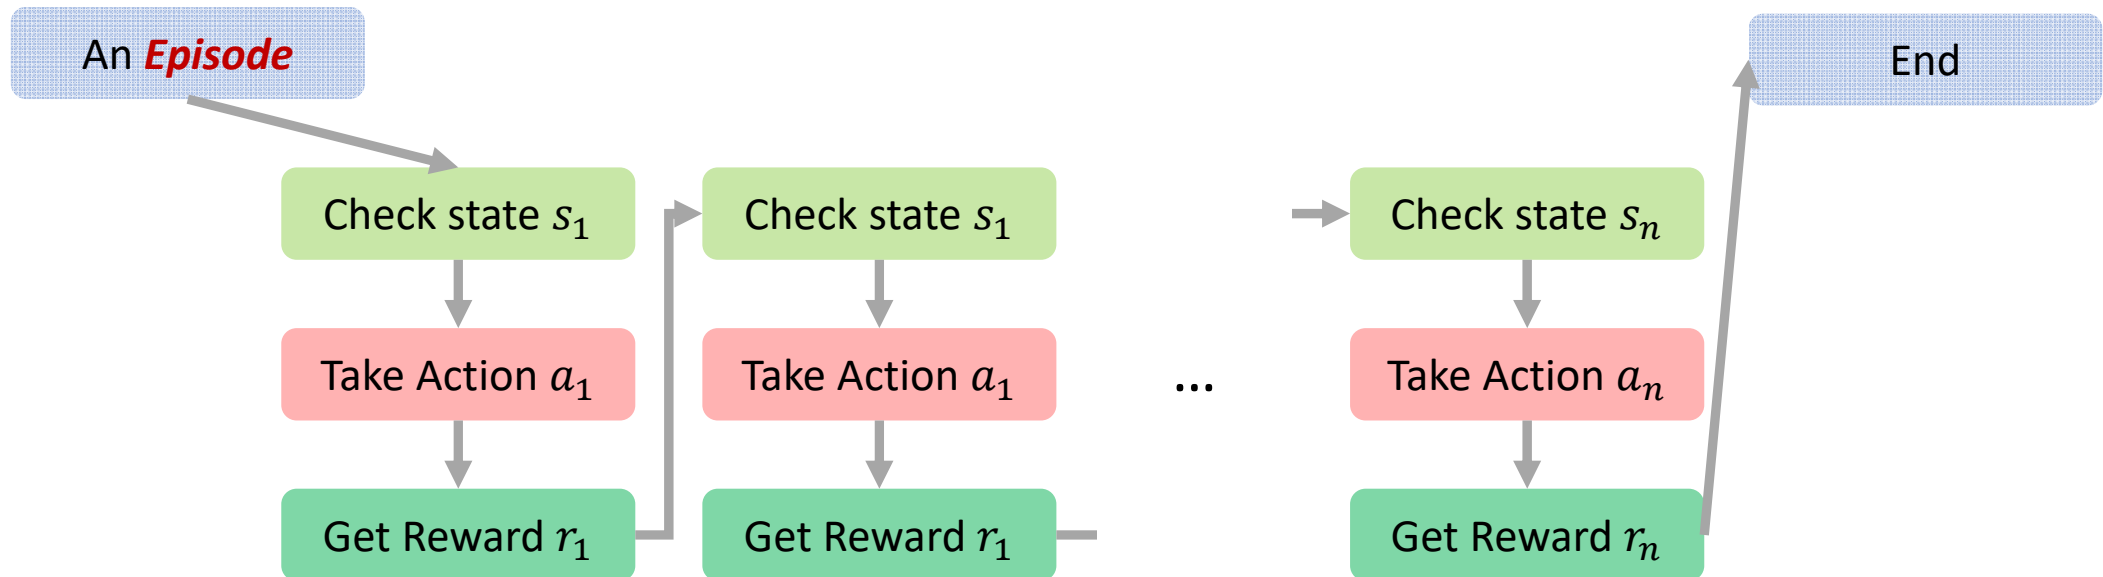
Take Action



Learning Process

$Data = \{(s_1, a_1), (s_2, a_2), (s_3, a_3), \dots, (s_n, a_n)\}$

$Reward = \{r_1, r_2, r_3, \dots, r_n\}$



Supervised vs Reinforcement



◆ For supervised Learning, we should have the dataset of *(state, action)* pairs ready from previous experience.

◆ Neglect about reward.

◆ Train a machine learning model.

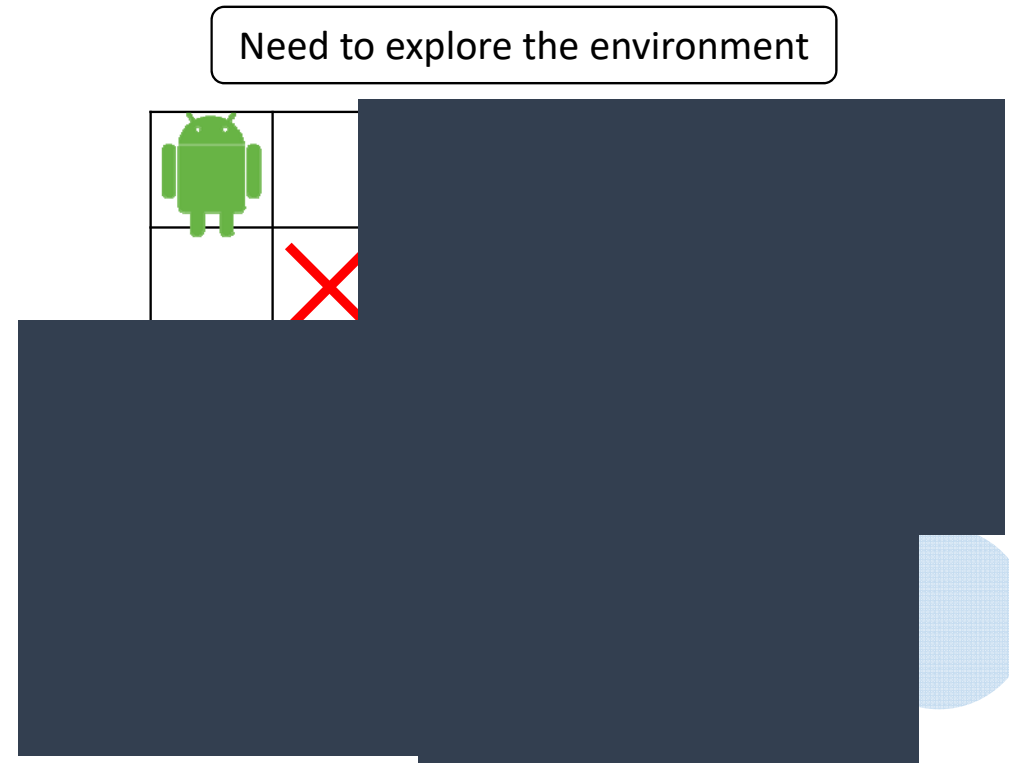
◆ Copy the experience.

$s_1 \rightarrow a_1$
 $s_2 \rightarrow a_2$
 $s_3 \rightarrow a_3$
.....



Supervised vs Reinforcement

- ◆ For reinforcement Learning, sometime the dataset is not ready.
- ◆ We need to explore during the learning process.
- ◆ We have to consider the total reward.





Value-based v.s. Policy-based

◆ Value-based: Learn a model to evaluate the “total reward” from the state.

◆ Q-Learning

◆ DQN



◆ Policy-based: Learn a model to decide “action” from the state.


◆ Policy gradient

Value-based v.s. Policy-based

Value-based



Agent



Action Up


Action Down

	Up	Down
State A	10	30
State B	30	21

Policy-based




Agent

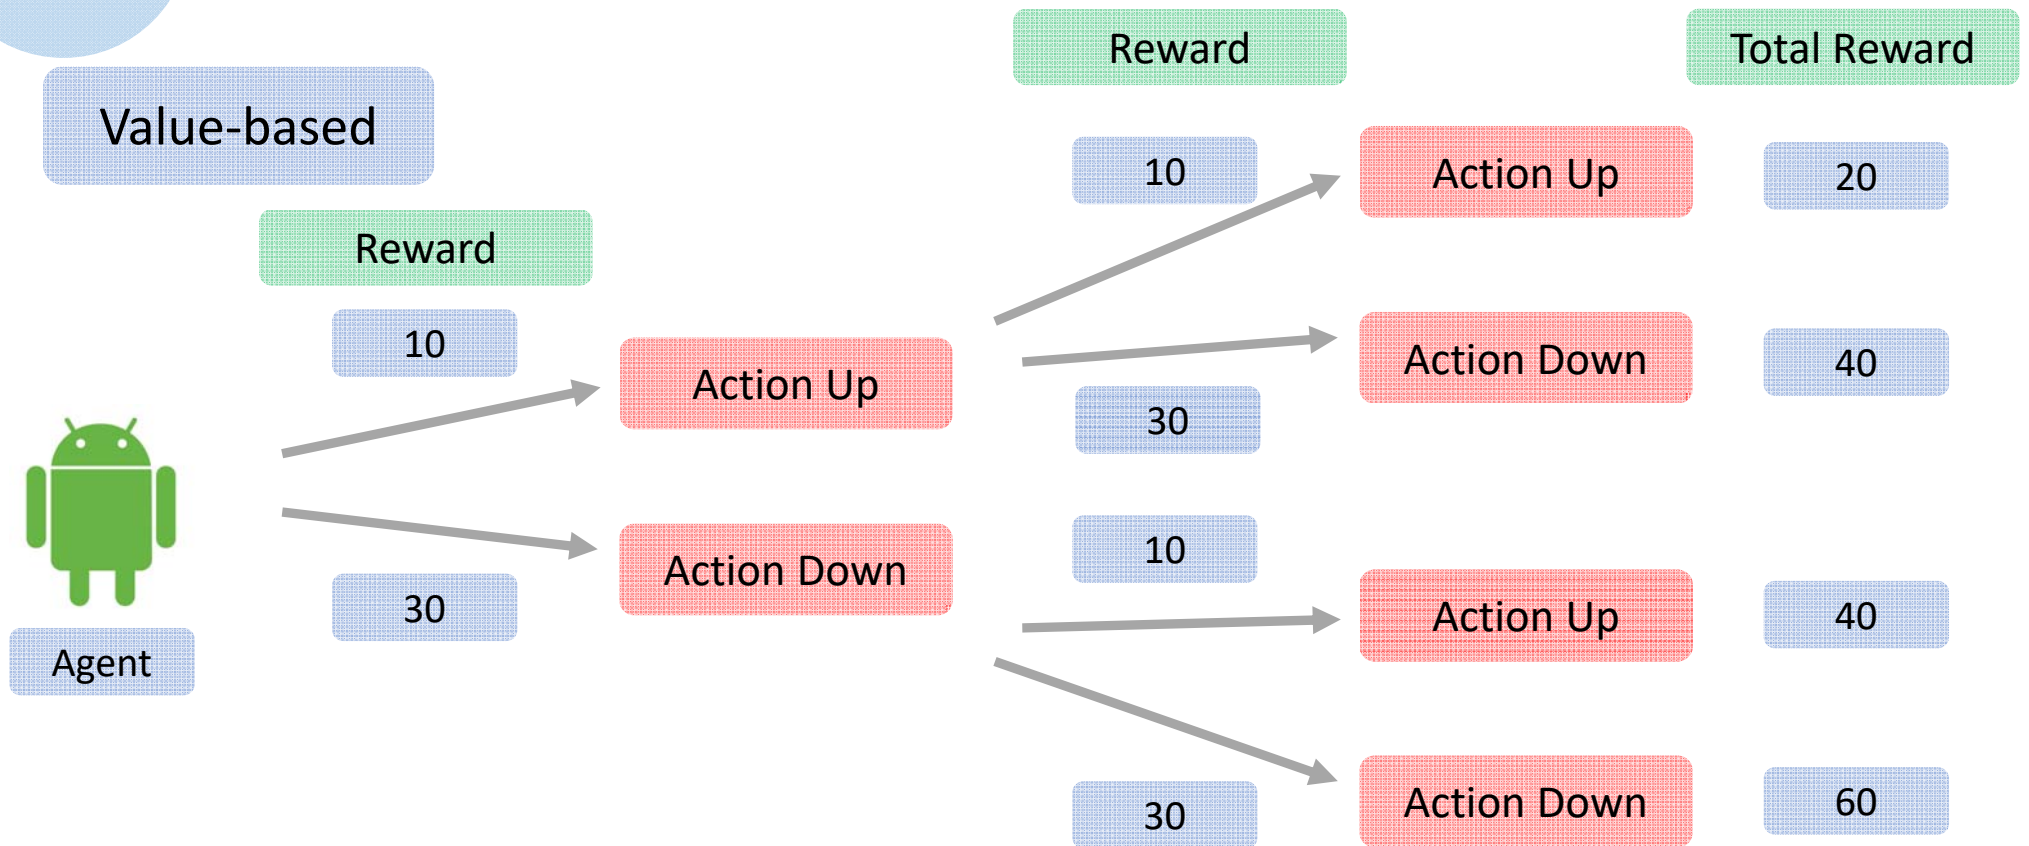


Action Up

Action Down



Value-based




Q-learning

- ◆ Keep update a so-called “**Q-table**”.
- ◆ Q-table represents the expected future total reward for each (*state*, *action*) pair.
- ◆ Q-table is built and updated through exploration.

Q-table

State	a_1	a_2	...	a_L
s_1				
s_2				
...				
s_N				



Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

Initialize s

Repeat (for each step of episode):

Choose a from s using policy derived from Q (e.g., ϵ -greedy)

Take action a , observe r, s'

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r_t + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

Estimate of optimal future reward

$s \leftarrow s'$;

until s is terminal

Q-table


State	a_1	a_2	...	a_L
s_1				
s_2				
...				
s_N				

α : learning rate

γ : discount factor

$$(1 - \alpha)Q(s, a) + \alpha [r_t + \gamma \max_{a'} Q(s', a')]$$

◆ After taking action a at state s and receiving the reward r_t , the expected reward of $Q(s, a)$ will be updated.



Choose the best action

Discount Factor in Q learning

◆ The discount factor γ determines the importance of future rewards.

$$Q(s_1, a) = (1 - \alpha)Q(s_1, a) + \alpha [r_t + \gamma \max_{a'} Q(s_1', a')]$$

$$\alpha = 1 \quad Q(s_1, a) = r_1 + \gamma Q(s_2, a) = r_1 + \gamma[r_2 + \gamma Q(s_3, a)]$$


$$Q(s_1, a) = r_1 + \gamma r_2 + \gamma^2 r_3 + \gamma^3 r_4 + \dots$$

$$\gamma = 1 \quad Q(s_1, a) = r_1 + r_2 + r_3 + r_4 + \dots$$

$$\gamma = 0 \quad Q(s_1, a) = r_1$$

Q-learning example

◆ We want to get the diamond.




Action	Reward
Nothing	-1
Trap	-100
Diamond	10

Game over

Initial Q-table

- ◆ We want to get a diamond.
- ◆ **X** represents the direction we can't take.



Q-table

Action	Reward
Nothing	-1
Trap	-100
Diamond	10


	Up	Down	Left	Right
1	X	0	X	0
2	X	0	0	0
3	X	0	0	X
4	0	0	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	0
8	0	X	0	0
9	0	X	0	X

Q-learning

Episode 1

Action	Reward
Nothing	-1
Trap	-100
Diamond	10

- ◆ We will find out the action that can achieve the expected maximum reward according to the Q-table.



	Up	Down	Left	Right
1	X	0	X	0
2	X	0	0	0
3	X	0	0	X
4	0	0	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	0
8	0	X	0	0
9	0	X	0	X


Q-learning

Episode 1

◆ Update $Q(1, \text{Right})$

$$\diamond (1 - \alpha)Q(1, \text{Right}) + \alpha [-1 + \gamma \max_{a'} Q(2, a')] = -1$$

Action	Reward
Nothing	-1
Trap	-100
Diamond	10



	Up	Down	Left	Right
1	X	0	X	-1
2	X	0	0	0
3	X	0	0	X
4	0	0	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	0
8	0	X	0	0
9	0	X	0	X

Update


Q-learning

Episode 1

Update $Q(2, \text{Down})$

$$(1 - \alpha)Q(2, \text{Down}) + \alpha [-100 + \gamma \max_{a'} Q(5, a')] = -100$$

Action	Reward
Nothing	-1
Trap	-100
Diamond	10



Reward = -100

Game over

	Up	Down	Left	Right
1	X	0	X	-1
2	X	-100	0	0
3	X	0	0	X
4	0	0	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	0
8	0	X	0	0
9	0	X	0	X


Update

Q-learning

Episode 2

Action	Reward
Nothing	-1
Trap	-100
Diamond	10

◆ We choose the action that we expect to get the most reward in the future.



	Up	Down	Left	Right
1	X	0	X	-1
2	X	-100	0	0
3	X	0	0	X
4	0	0	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	0
8	0	X	0	0
9	0	X	0	X


Q-learning

Episode 2

Update $Q(1, \text{Down})$

$$(1 - \alpha)Q(1, \text{Down}) + \alpha [-1 + \gamma \max_{a'} Q(4, a')] = -1$$

Action	Reward
Nothing	-1
Trap	-100
Diamond	10



	Up	Down	Left	Right
1	X	-1	X	-1
2	X	-100	0	0
3	X	0	0	X
4	0	0	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	0
8	0	X	0	0
9	0	X	0	X

Update


Q-learning

Episode 2

Action	Reward
Nothing	-1
Trap	-100
Diamond	10

Update $Q(4, \text{Down})$

$$(1 - \alpha)Q(4, \text{Down}) + \alpha [-1 + \gamma \max_{a'} Q(7, a')] = -1$$



	Up	Down	Left	Right
1	X	-1	X	-1
2	X	-100	0	0
3	X	0	0	X
4	0	-1	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	0
8	0	X	0	0
9	0	X	0	X

Update


Q-learning

Episode 2

Action	Reward
Nothing	-1
Trap	-100
Diamond	10

◆ Update $Q(8, \text{Right})$

◆ $(1 - \alpha)Q(8, \text{Right}) + \alpha [10 + \gamma \max_{a'} Q(9, a')] = 10$



	Up	Down	Left	Right
1	X	0	X	-1
2	X	-100	0	0
3	X	0	0	X
4	0	-1	X	0
5	0	0	0	0
6	0	0	0	X
7	0	X	X	-1
8	0	X	0	10
9	0	X	0	X


Update

Q-learning

Episode 400

Action	Reward
Nothing	-1
Trap	-100
Diamond	10

◆ After many episodes, we will have an optimized Q-table.



$$\alpha = 1$$

$$\gamma = 1$$

$$\varepsilon = 1$$

$$\alpha = 0.9$$

$$\gamma = 0.95$$

$$\varepsilon = 0.9$$


	Up	Down	Left	Right
1	X	5.72	4.44	5.72
2	X	-100	4.44	7.07
3	X	8.5	5.72	X
4	4.44	7.07	X	-100
5	0	0	0	0
6	0	10	-100	X
7	5.72	X	X	8.5
8	0	X	7.07	10
9	0	0	0	0

Q-learning

Episode 400

Action	Reward
Nothing	-1
Trap	-100
Diamond	10

◆ We can find the best path according to Q-table.



	Up	Down	Left	Right
1	X	7	X	7
2	X	-100	6	8
3	X	9	7	X
4	6	8	X	-100
5	0	0	0	0
6	0	10	-100	X
7	7	X	X	9
8	0	X	8	10
9	0	0	0	0



Exploration and Exploitation

◆ Sometime we need to explore new paths for learning better.

◆ Exploration

◆ Try different actions even if you don't get the best reward.


◆ Exploitation

◆ Choose the best action that gets the highest reward.

Exploration and Exploitation


Exploitation

The path we chose to explore



Exploration


Try to explore new ways



Exploration and Exploitation

Exploration

Sometimes we can get better results(Shorter path).




In Q-learning, this means that we sometimes choose actions that currently will not receive the most expected value.

	Up	Down	Left	Right
1	X	-20	X	10
2	X	-45	-53	-30
3	X	-40	20	X
...
n

Maximum expected reward

Epsilon Greedy

- ◆ Assume we just use 'Greedy' policy.
- ◆ First time we choose 'Right' in 'A'



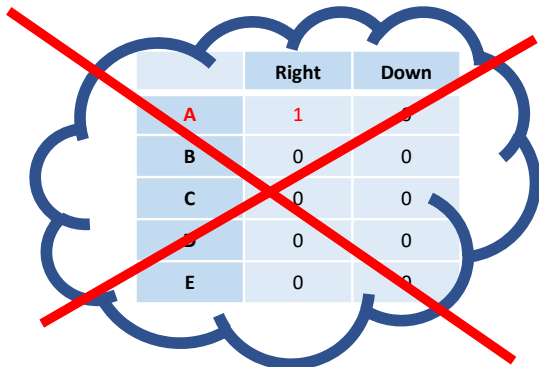
- We will never choose action 'Down' at state A.

Epoch = 0

state	Right	Down
A	1	0
B	0	0
C	0	0
D	0	0
E	0	0

Epsilon Greedy


◆ Sometime we choose action randomly in order to explore all possibilities



	Right	Down
A	1	0
B	0	0
C	0	0
D	0	0
E	0	0




Down



Q-table explosion

- ◆ If the number of *states/actions* is infinite or very large, Q-table becomes infeasible.


Infinite



Velocity



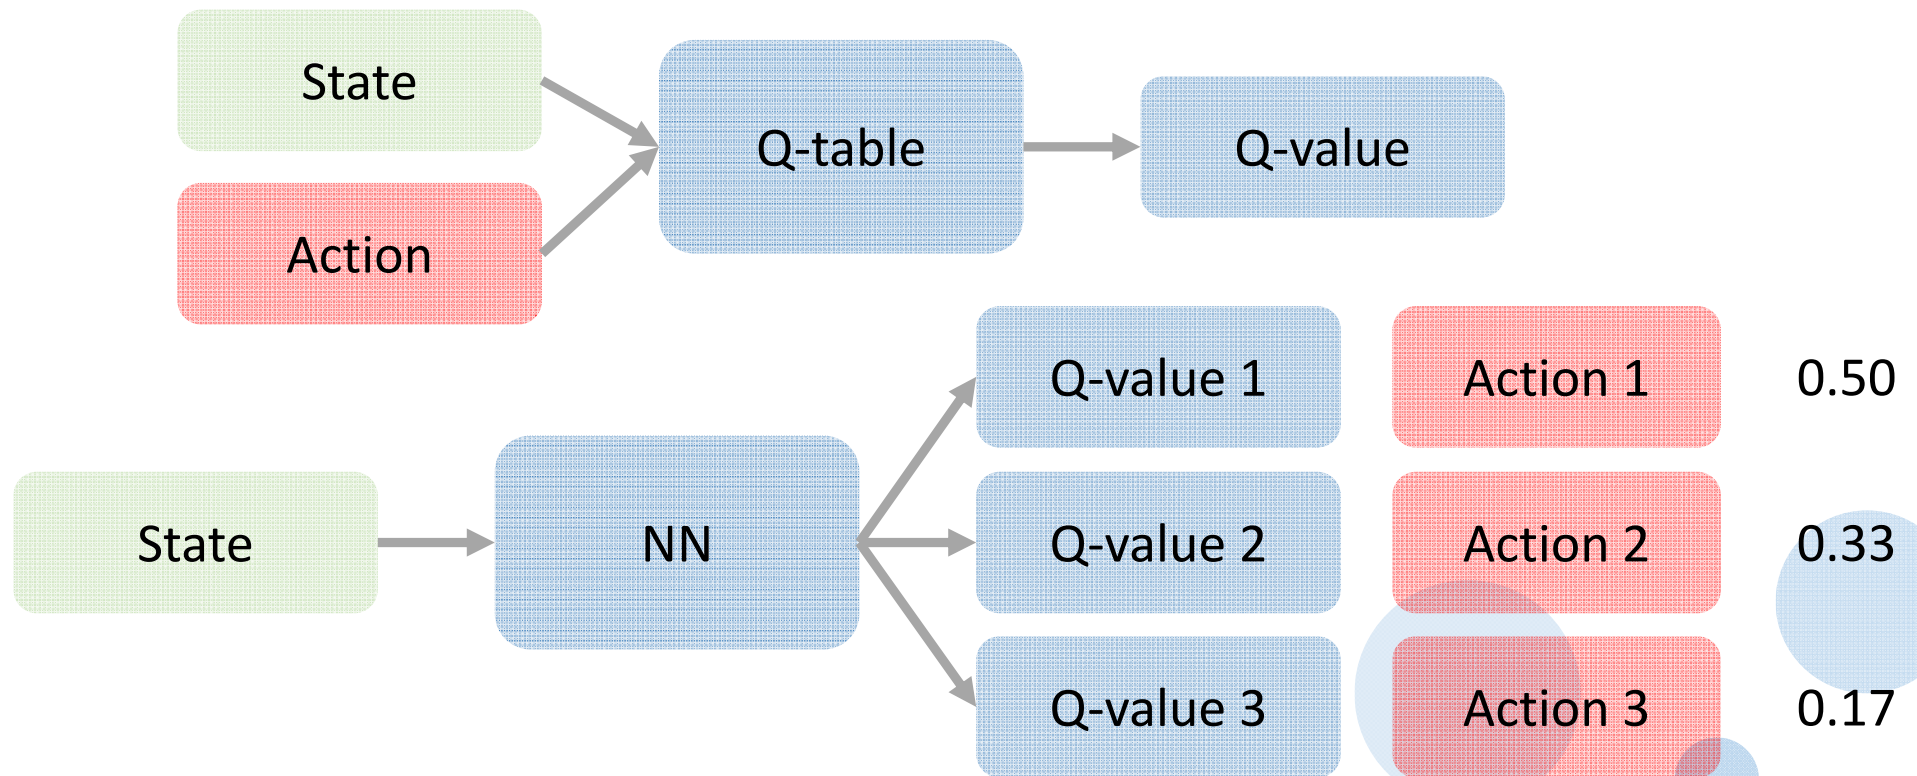
Large



- ◆ *How many possible states in GO?*

Deep Q Network(DQN)

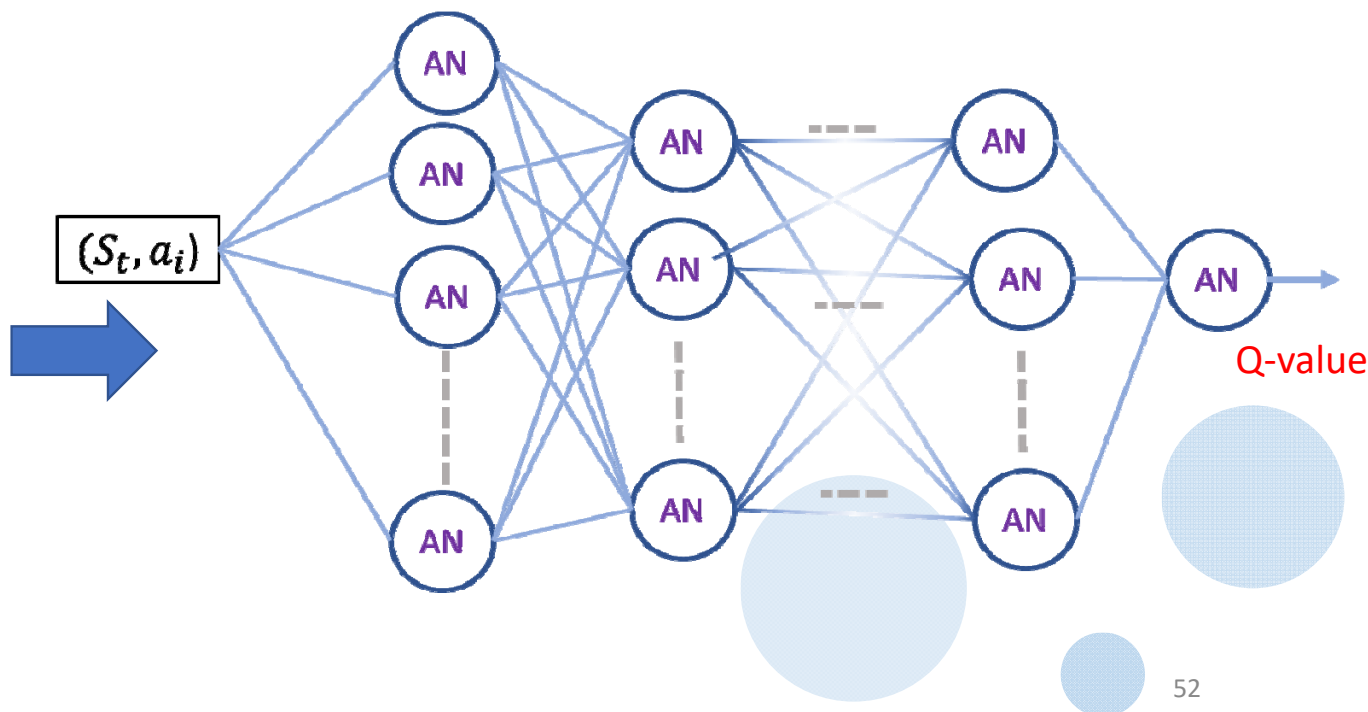
- ◆ Generate Q-value for each (*state*, *action*) pair.



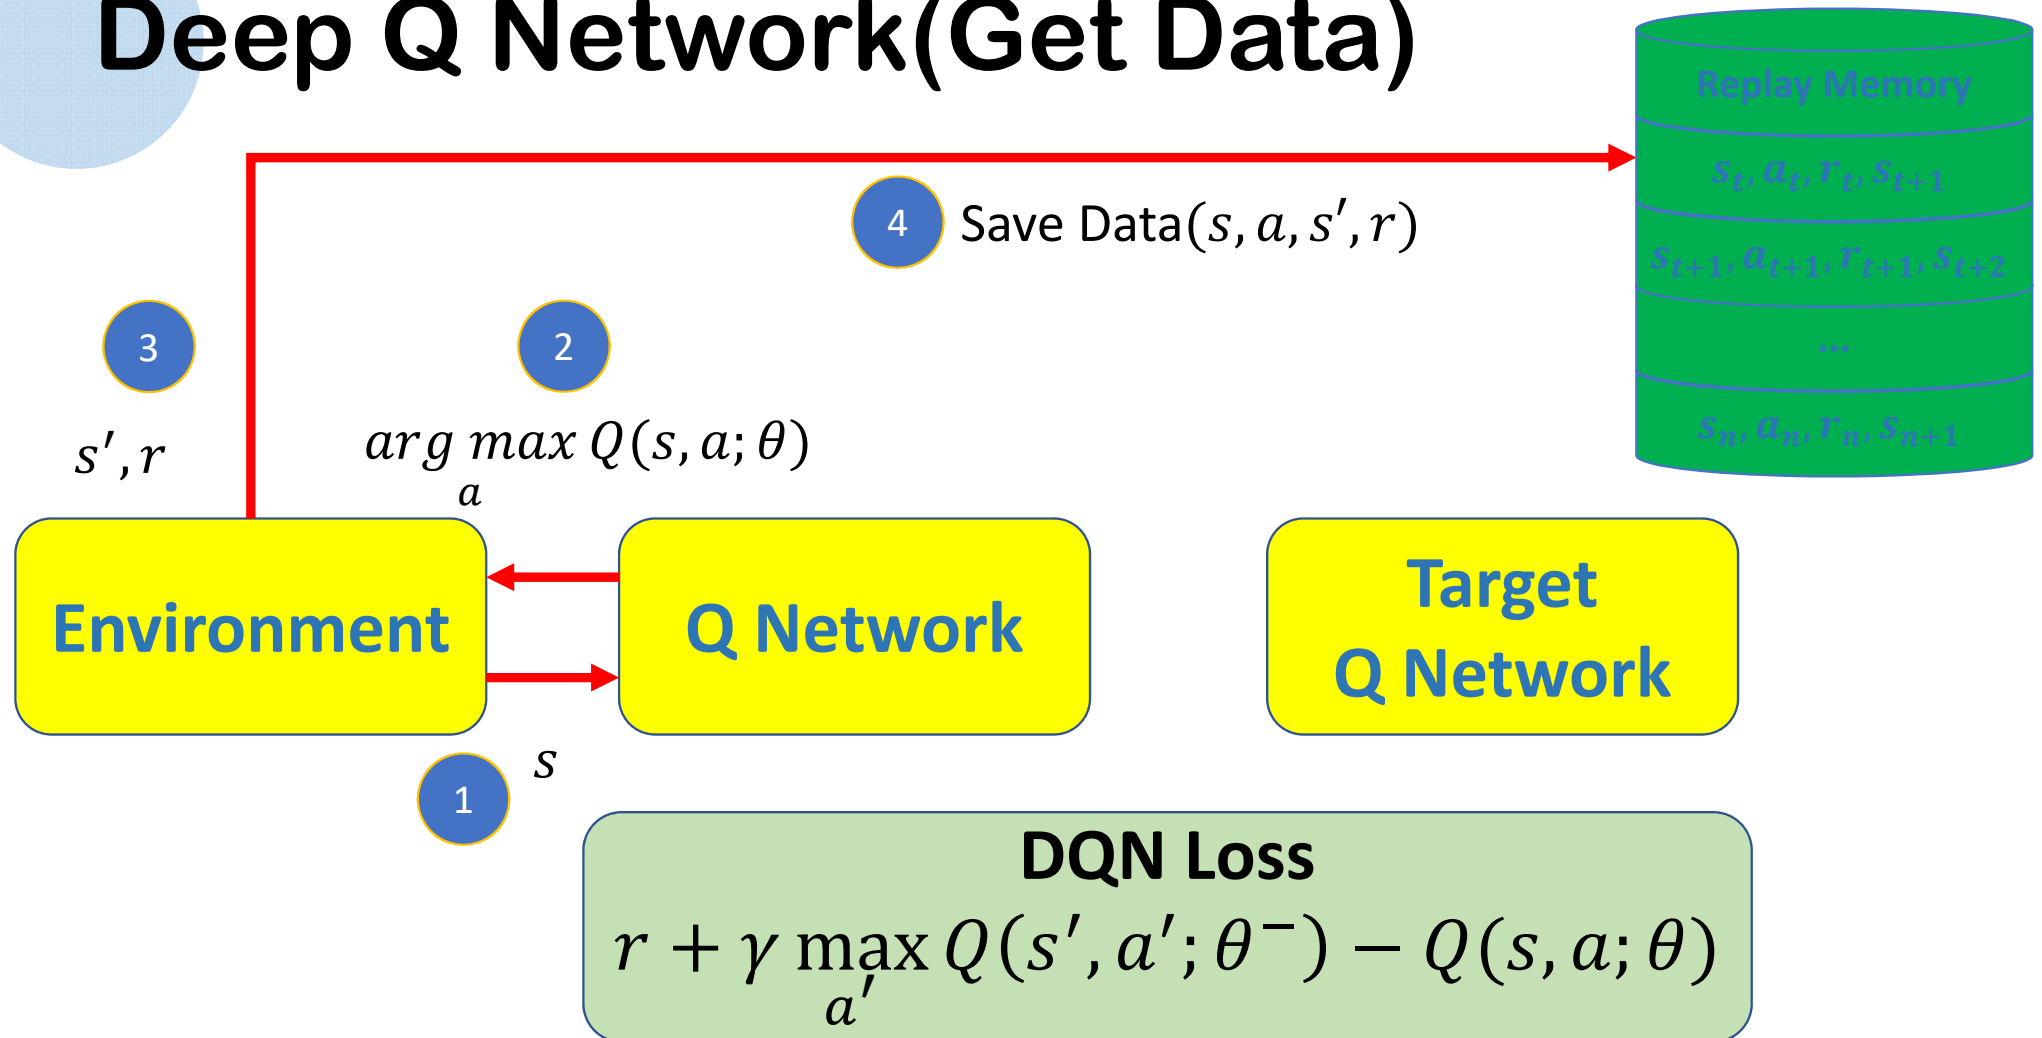
Build Q-table Model

- ◆ When # of *states* or *actions* is too big to record, we can use **Neural Network** to replace it.

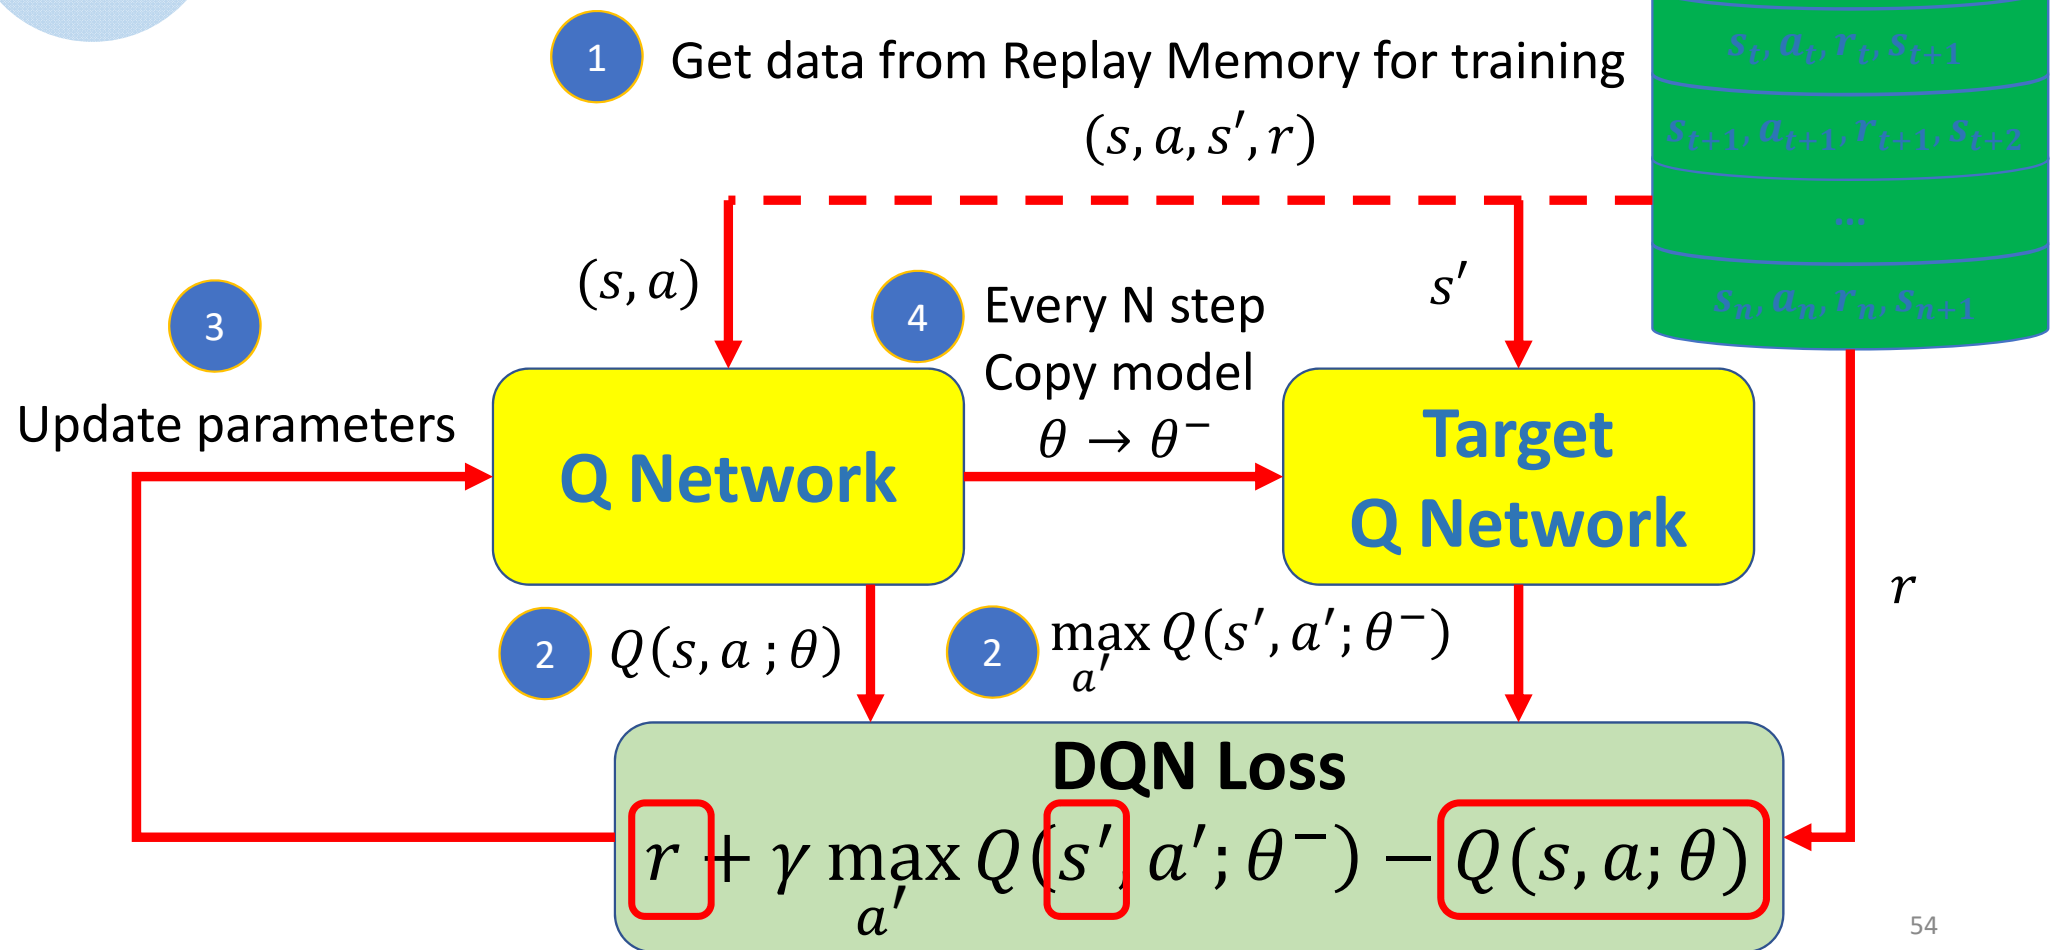
	Right	Down
S1	1	0
S2	0	0
S3	0	0
S4	0	0
S5	0	0
...



Deep Q Network(Get Data)



Deep Q Network(Training)



Fixed Q-targets

- ◆ Use two Network, **Q Network** and **Target Q Network**.
- ◆ Every N step, we will copy the **Q Network** to the **Target Q Network**.
- ◆ Target Q Network is the **old** Q Network

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r_t + \gamma \max_{a'} \underbrace{Q(s', a')}_{\text{Target Q Network}} - Q(s, a)]$$

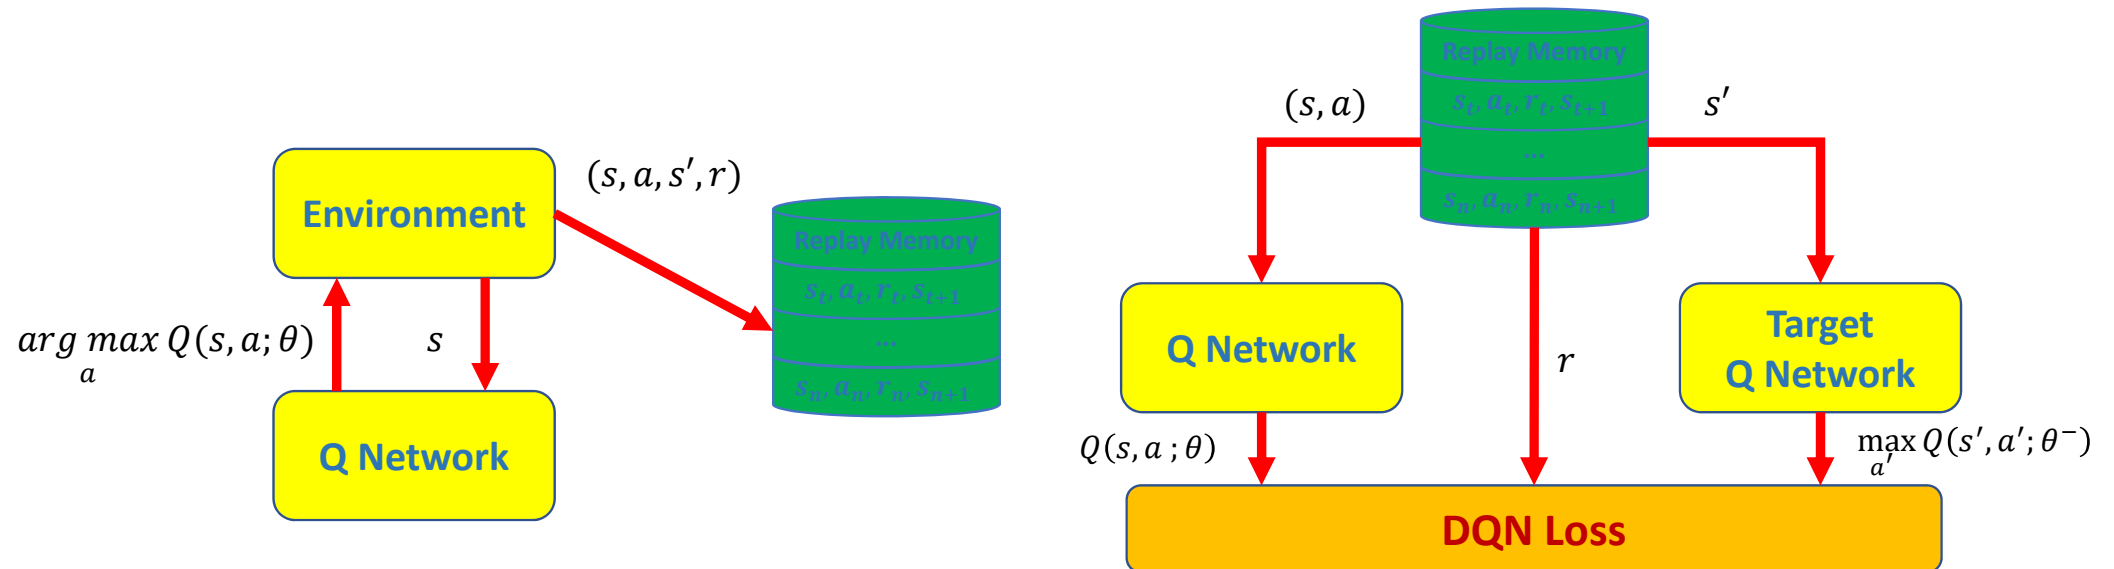
Q Network \longrightarrow **Target Q Network**

Every N step, copy model to Target Q Network

Experience Replay

Add data

Use data



Double Deep Q Network

- ◆ Q-learning and Deep Q Learning tends to overestimate q-values.
- ◆ Because we take the max over all actions.



$$Q(s, a) \leftarrow Q(s, a) + \alpha [r_t + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

- ◆ If all values were equally overestimated this would be no problem.
- ◆ But if the overestimations are not uniform, this might slow down learning.

Overestimated

◆ $X1, X2$ is sequence

◆ $E(\max(X1, X2)) \geq \max(E(X1), E(X2))$



Double Q-learning

- ◆ Use **two Q network**, one of them **evaluate**, another **choose an action**.
- ◆ Randomly exchange behavior of two Q networks.

```
Initialize  $Q_1(s, a)$  and  $Q_2(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily
Initialize  $Q_1(\text{terminal-state}, \cdot) = Q_2(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
  Initialize  $S$ 
  Repeat (for each step of episode):
    Choose  $A$  from  $S$  using policy derived from  $Q_1$  and  $Q_2$  (e.g.,  $\epsilon$ -greedy in  $Q_1 + Q_2$ )
    Take action  $A$ , observe  $R, S'$ 
    With 0.5 probability:
       $Q_1(S, A) \leftarrow Q_1(S, A) + \alpha(R + \gamma Q_2(S', \arg\max_a Q_1(S', a)) - Q_1(S, A))$ 
    else:
       $Q_2(S, A) \leftarrow Q_2(S, A) + \alpha(R + \gamma Q_1(S', \arg\max_a Q_2(S', a)) - Q_2(S, A))$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```

Double Q-learning

- ◆ Use **two Q network**, one of them **evaluate**, another **choose an action**.
- ◆ Randomly exchange behavior of two Q networks.

Episode 1

Choose an action

Q Network(A)

Evaluate value
(target network)

Q Network(B)

Double Q-learning

- ◆ Use **two Q network**, one of them **evaluate**, another **choose an action**.
- ◆ Randomly exchange behavior of two Q networks.

Episode 2

Evaluate value
(target network)

Choose an action

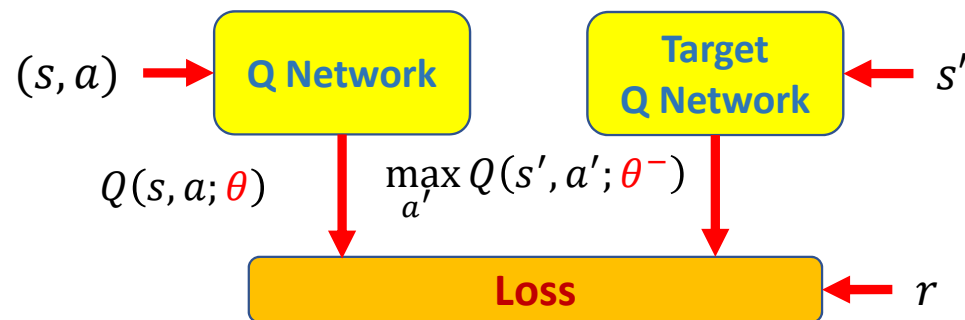
Change Behavior

Q Network(A)

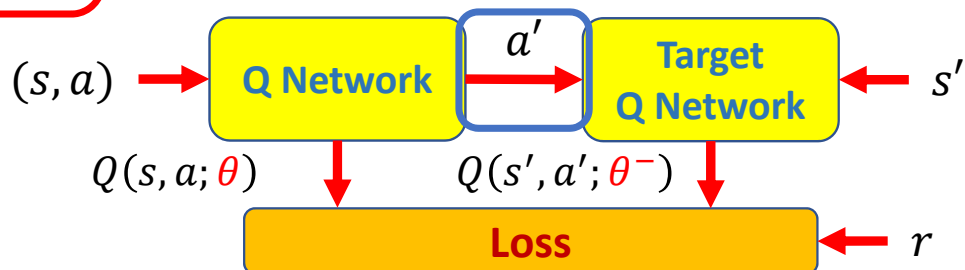
Q Network(B)

Double Deep Q Network

$$Y_t^Q = r_t + \gamma \max_{a'} Q(s', a'; \theta^-)$$




$$Y_t^{Double-Q} = r_t + \gamma Q\left(s', \arg \max_{a'} Q(s', a'; \theta); \theta^-\right)$$



Policy Gradient


◆ Similar to multi-class classification problem.



Policy gradient

◆ If we want to learn how to walk.....

Episode 1




Fail -100

Punish these actions, and try to avoid these actions

Policy gradient

◆ If we want to learn how to walk


Episode 3



Reward these actions, and choose these actions more often.

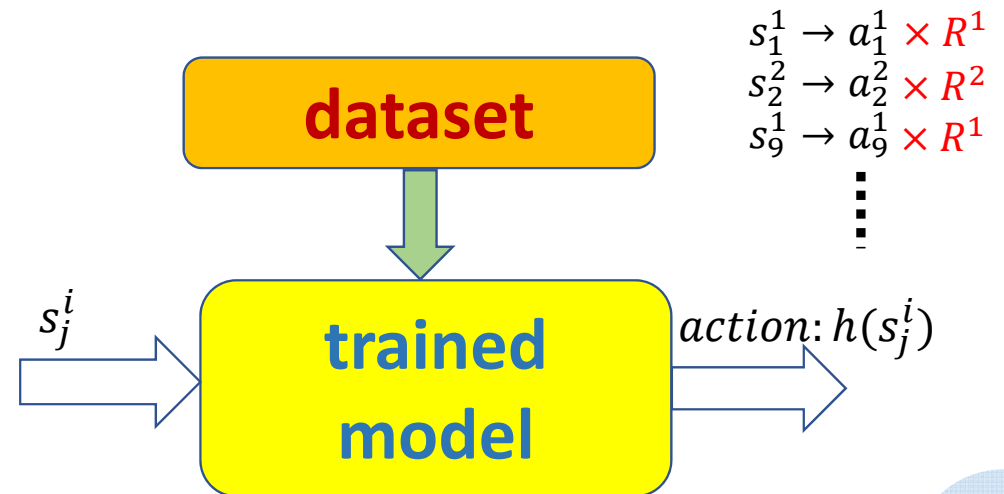
Policy Gradient Training

Supervised Learning



Episode 1: $[(s_1^1 \rightarrow a_1^1), (s_2^1 \rightarrow a_2^1), \dots, (s_{100}^1 \rightarrow a_{100}^1)]$
 Episode 2: $[(s_1^2 \rightarrow a_1^2), (s_2^2 \rightarrow a_2^2), \dots, (s_{70}^2 \rightarrow a_{70}^2)]$

Policy Gradient Training



Total reward: $R^1 = r_1^1 + r_2^1 \dots + r_{100}^1$
 Total reward: $R^2 = r_1^2 + r_2^2 \dots + r_{70}^2$



Monte Carlo & Temporal-Difference

- ◆ Monte Carlo Policy Gradient

- ◆ Update every episode.

- ◆ Temporal-Difference


- ◆ Update every step.

Monte Carlo

◆ Monte Carlo Policy Gradient

◆ Update every episode.

Episode 1




$$episode(1) = \{(s_1, a_1, r_1), (s_2, a_2, r_2), (s_3, a_3, r_3), \dots, (s_n, a_n, r_n)\}$$

Update 1 time

Temporal-Difference

◆ Temporal-Difference

◆ Update every step.




$$episode(1) = \{(s_1, a_1, r_1), (s_2, a_2, r_2), (s_3, a_3, r_3), \dots, (s_n, a_n, r_n)\}$$

Update 5 times

Temporal-Difference

◆ Temporal-Difference

◆ Update every step.




$$episode(1) = \{(s_1, a_1, r_1), (s_2, a_2, r_2), (s_3, a_3, r_3), \dots, (s_n, a_n, r_n)\}$$

Update 2 times

Temporal-Difference

◆ Temporal-Difference

◆ Update every step.




$$episode(1) = \{(s_1, a_1, r_1), (s_2, a_2, r_2), (s_3, a_3, r_3), \dots, (s_n, a_n, r_n)\}$$

Temporal-Difference

◆ Temporal-Difference

◆ Update every step.




$$episode(1) = \{(s_1, a_1, r_1), (s_2, a_2, r_2), (s_3, a_3, r_3), \dots, (s_n, a_n, r_n)\}$$

Update 4 times

Temporal-Difference

◆ Temporal-Difference

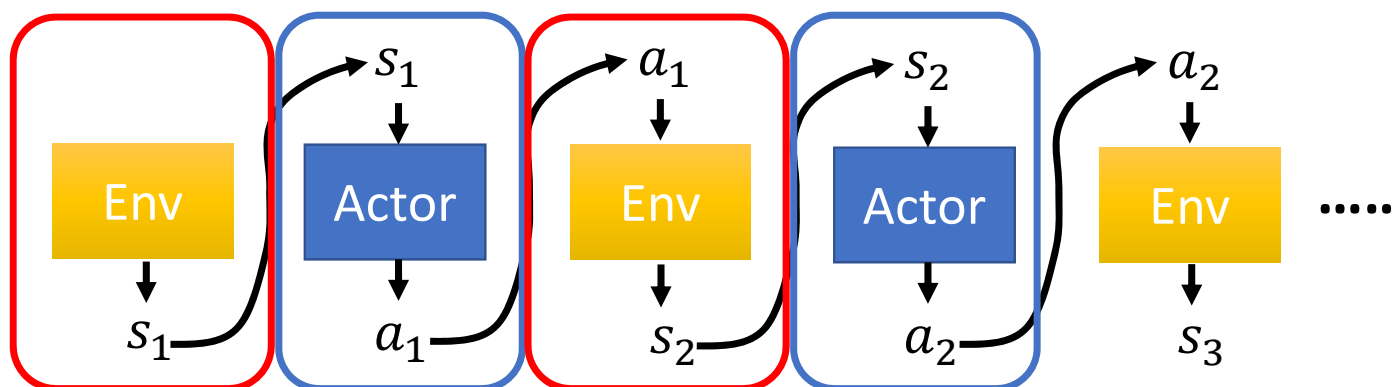
◆ Update every step.



$$episode(1) = \{(s_1, a_1, r_1), (s_2, a_2, r_2), (s_3, a_3, r_3), \dots, (s_n, a_n, r_n)\}$$

Update 5 times

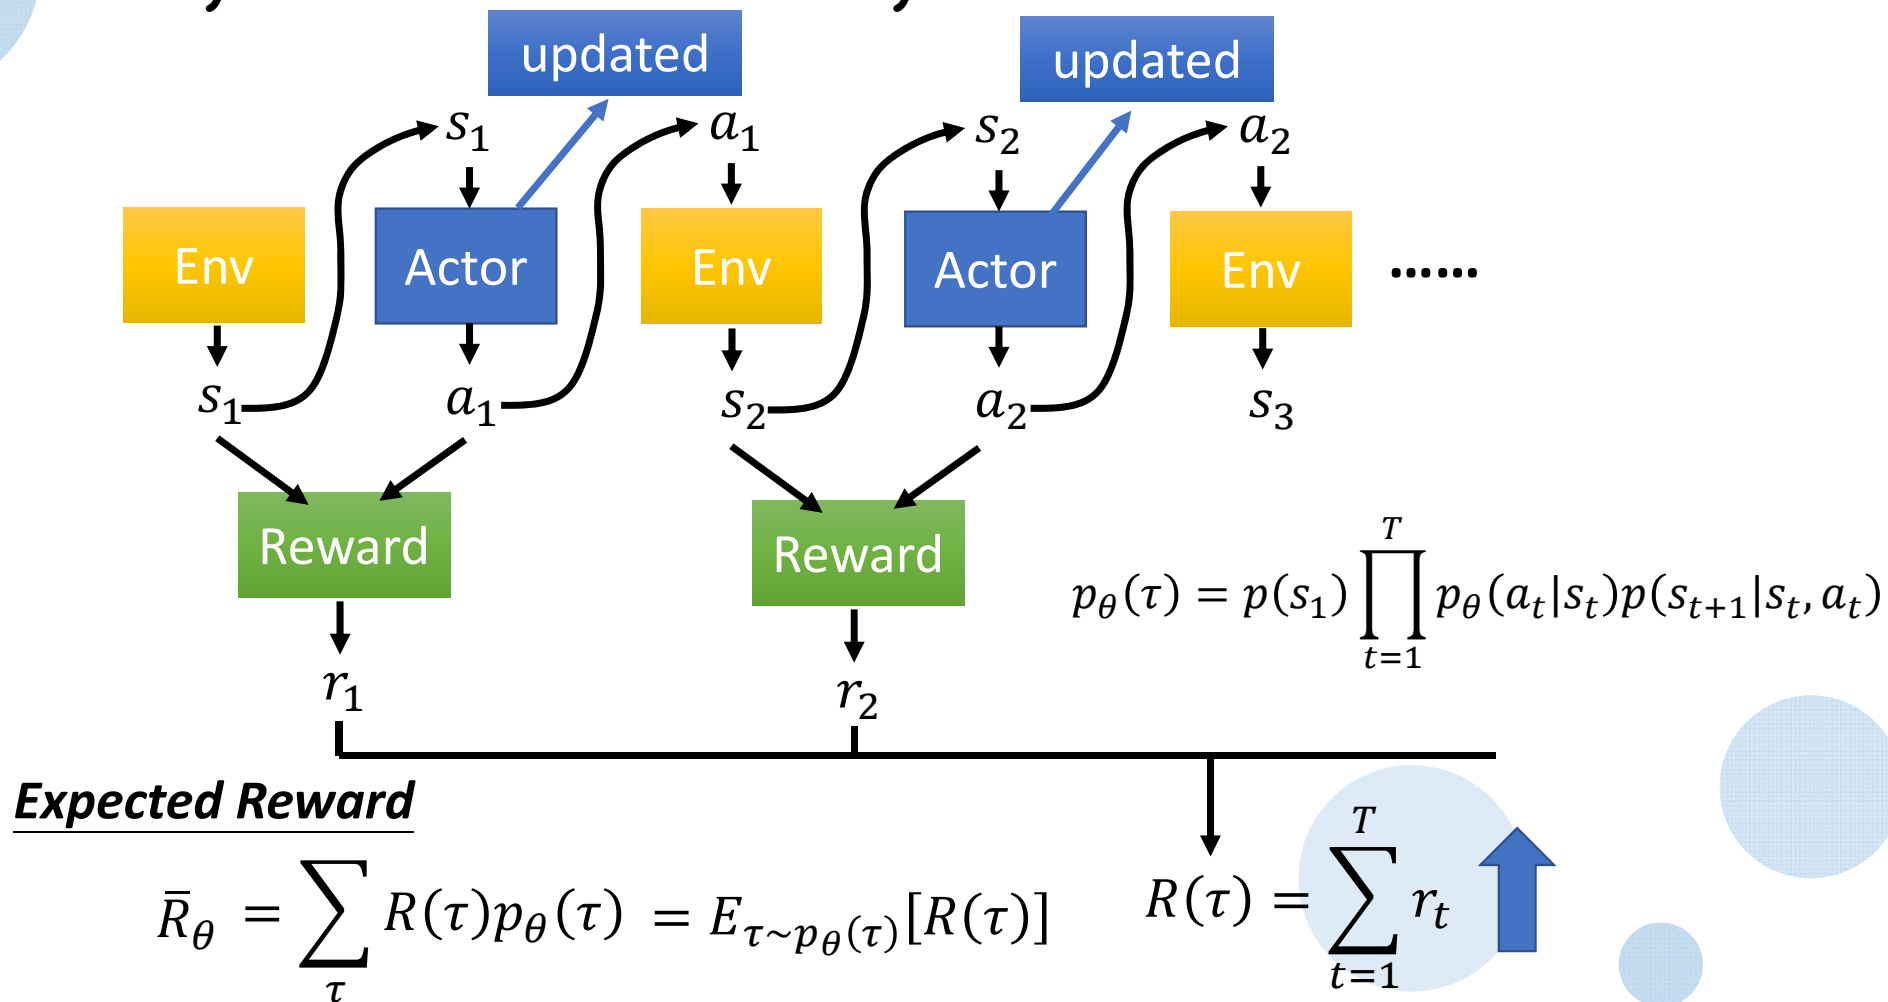
Actor, Environment, Reward



Trajectory $\tau = \{s_1, a_1, s_2, a_2, \dots, s_T, a_T\}$

$$\begin{aligned} p_{\theta}(\tau) &= p(s_1) p_{\theta}(a_1 | s_1) p(s_2 | s_1, a_1) p_{\theta}(a_2 | s_2) p(s_3 | s_2, a_2) \dots \\ &= p(s_1) \prod_{t=1}^T p_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t) \end{aligned}$$

Actor, Environment, Reward



Policy Gradient

$$\bar{R}_\theta = \sum_{\tau} R(\tau) p_\theta(\tau) \quad \nabla \bar{R}_\theta = ?$$

$$\nabla \bar{R}_\theta = \sum_{\tau} R(\tau) \nabla p_\theta(\tau) = \sum_{\tau} R(\tau) p_\theta(\tau) \frac{\nabla p_\theta(\tau)}{p_\theta(\tau)}$$

$$\nabla f(x) = f(x) \nabla \log f(x)$$

$R(\tau)$ do not have to be differentiable

It can even be a black box.

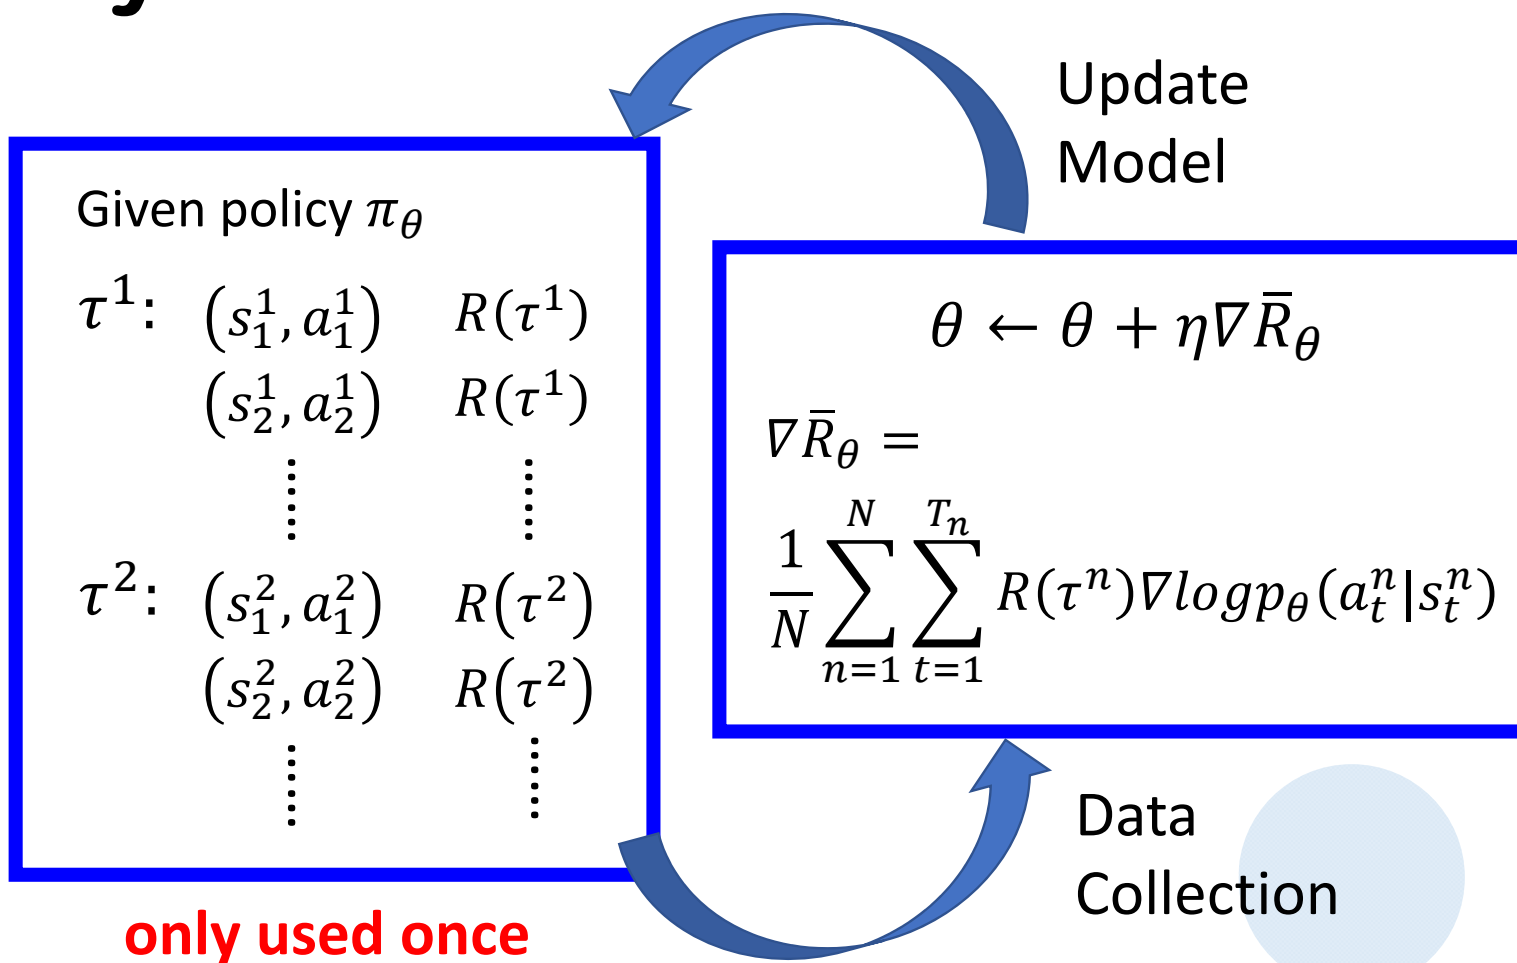
$$= \sum_{\tau} R(\tau) p_\theta(\tau) \nabla \log p_\theta(\tau)$$

$$= E_{\tau \sim p_\theta(\tau)} [R(\tau) \nabla \log p_\theta(\tau)] \approx \frac{1}{N} \sum_{n=1}^N R(\tau^n) \nabla \log p_\theta(\tau^n)$$

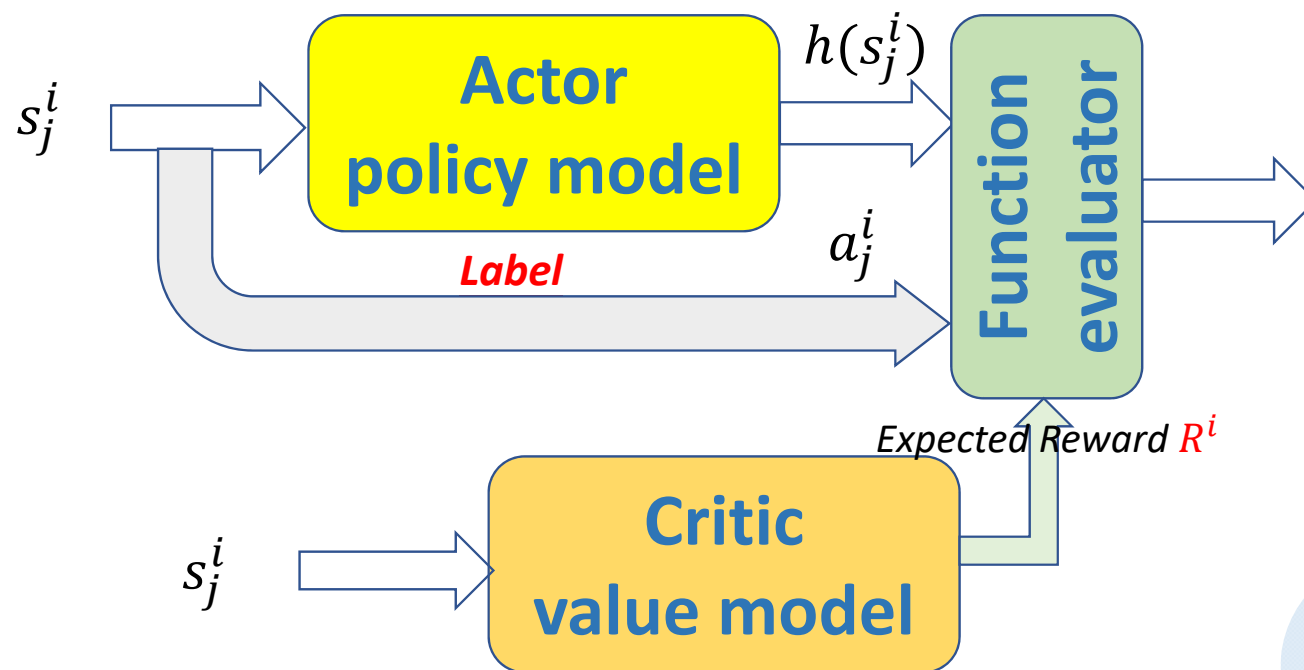
$$= \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} R(\tau^n) \nabla \log p_\theta(a_t^n | s_t^n)$$

Policy Gradient

$$\nabla \bar{R}_\theta = E_{\tau \sim p_\theta(\tau)} [R(\tau) \nabla \log p_\theta(\tau)]$$

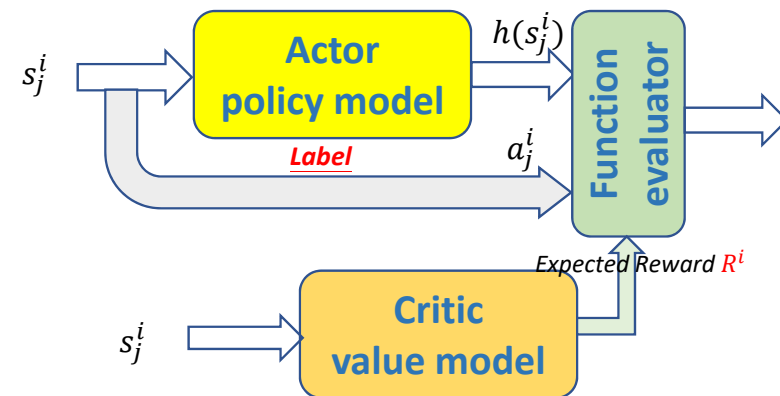


Actor-Critic



Actor-Critic

- ◆ Actor : Policy base
- ◆ Critic : Value base
- ◆ Actor selects the action, Critic evaluates the quality of the selected action




Actor-Critic

Actor

left

right

straight



Agent

Critic

None


Actor-Critic

Actor

left

right

straight



Agent

Critic

-10


Actor-Critic

Actor

left

right

straight



Agent

Critic

-20


Actor-Critic

Actor

left

right

straight



Agent

Critic

10

Actor-Critic

Actor

left

right

straight



Agent

Critic

10

Actor-Critic

◆ Policy Gradient

$$\nabla \bar{R}_\theta \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\sum_{t'=t}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_\theta(a_t^n | s_t^n)$$

G_t^n : obtained via interaction

Actor-Critic

- ◆ Policy Gradient
- ◆ Use critic to calculate possible future rewards

$$\nabla \bar{R}_\theta \approx \frac{1}{N} \sum_{n=1}^N \sum_{t=1}^{T_n} \left(\sum_{t'=t}^{T_n} \gamma^{t'-t} r_{t'}^n - b \right) \nabla \log p_\theta(a_t^n | s_t^n)$$

G_t^n : obtained via interaction

$$E[G_t^n] = Q^{\pi_\theta}(s_t^n, a_t^n)$$

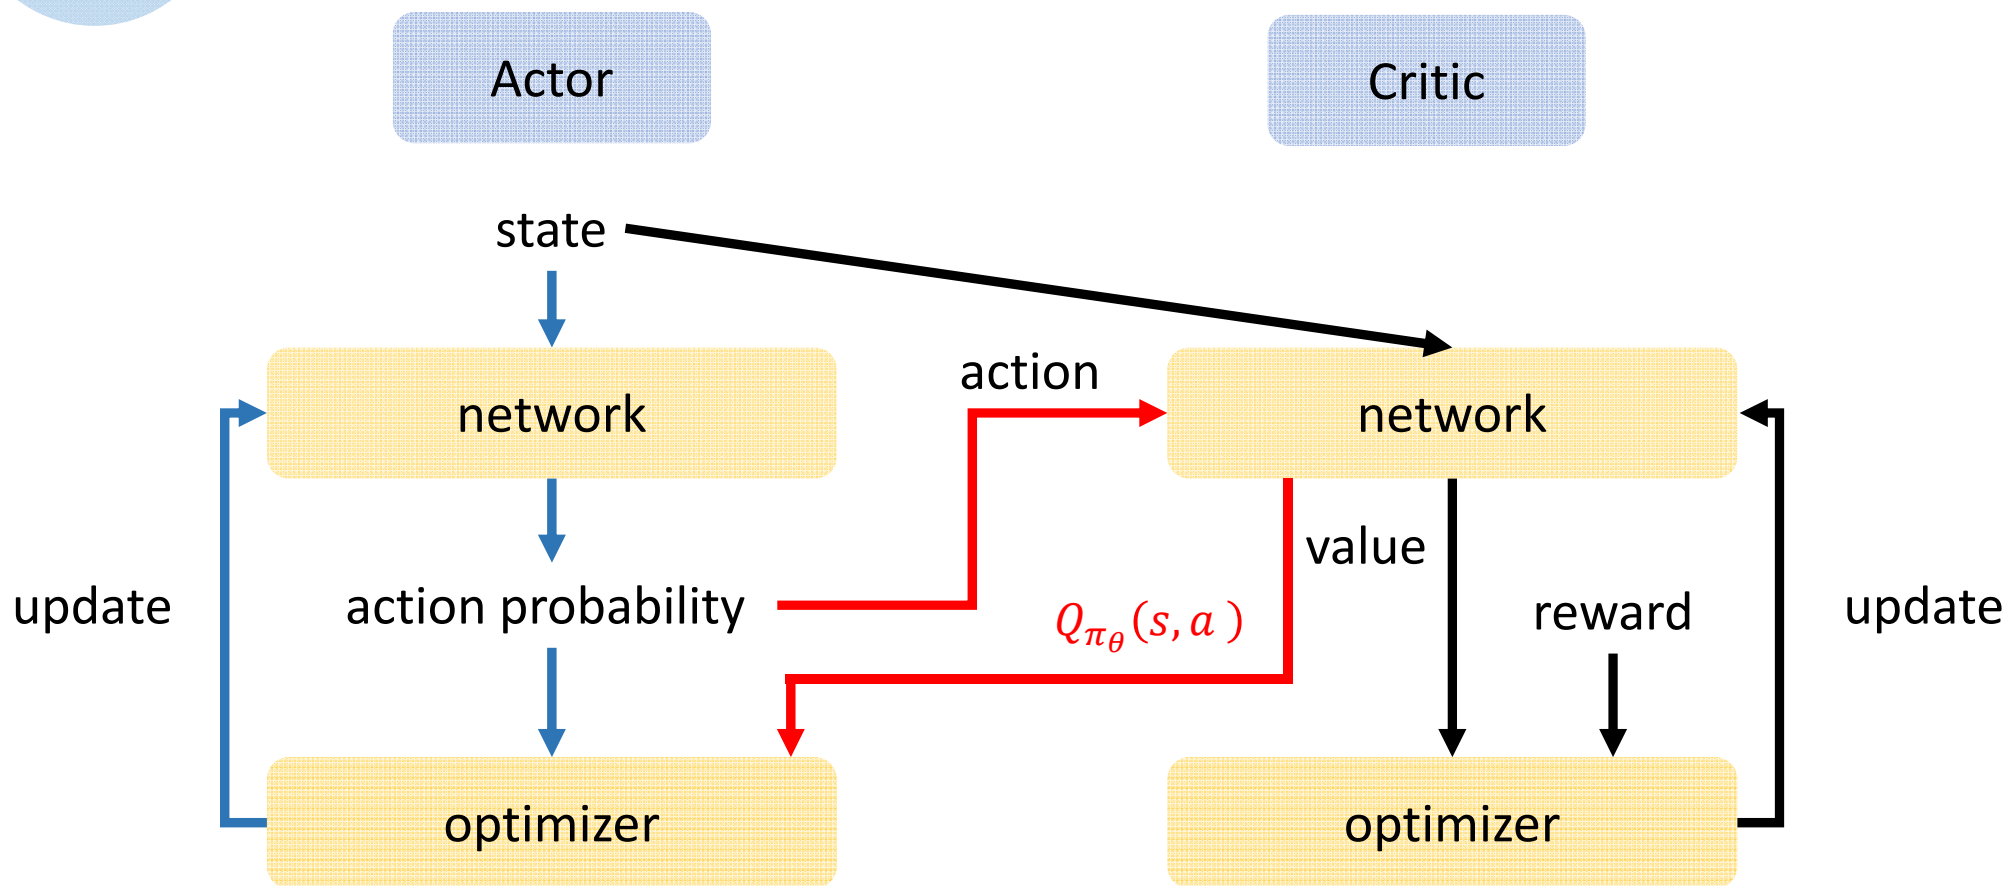
Actor-Critic

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) G_t] \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q^w(s, a)] \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^w(s, a)] \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \delta]\end{aligned}$$

REINFORCE
Q Actor-Critic
Advantage Actor-Critic
TD Actor-Critic

Actor-Critic

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q_{\pi_{\theta}}(s, a)]$$

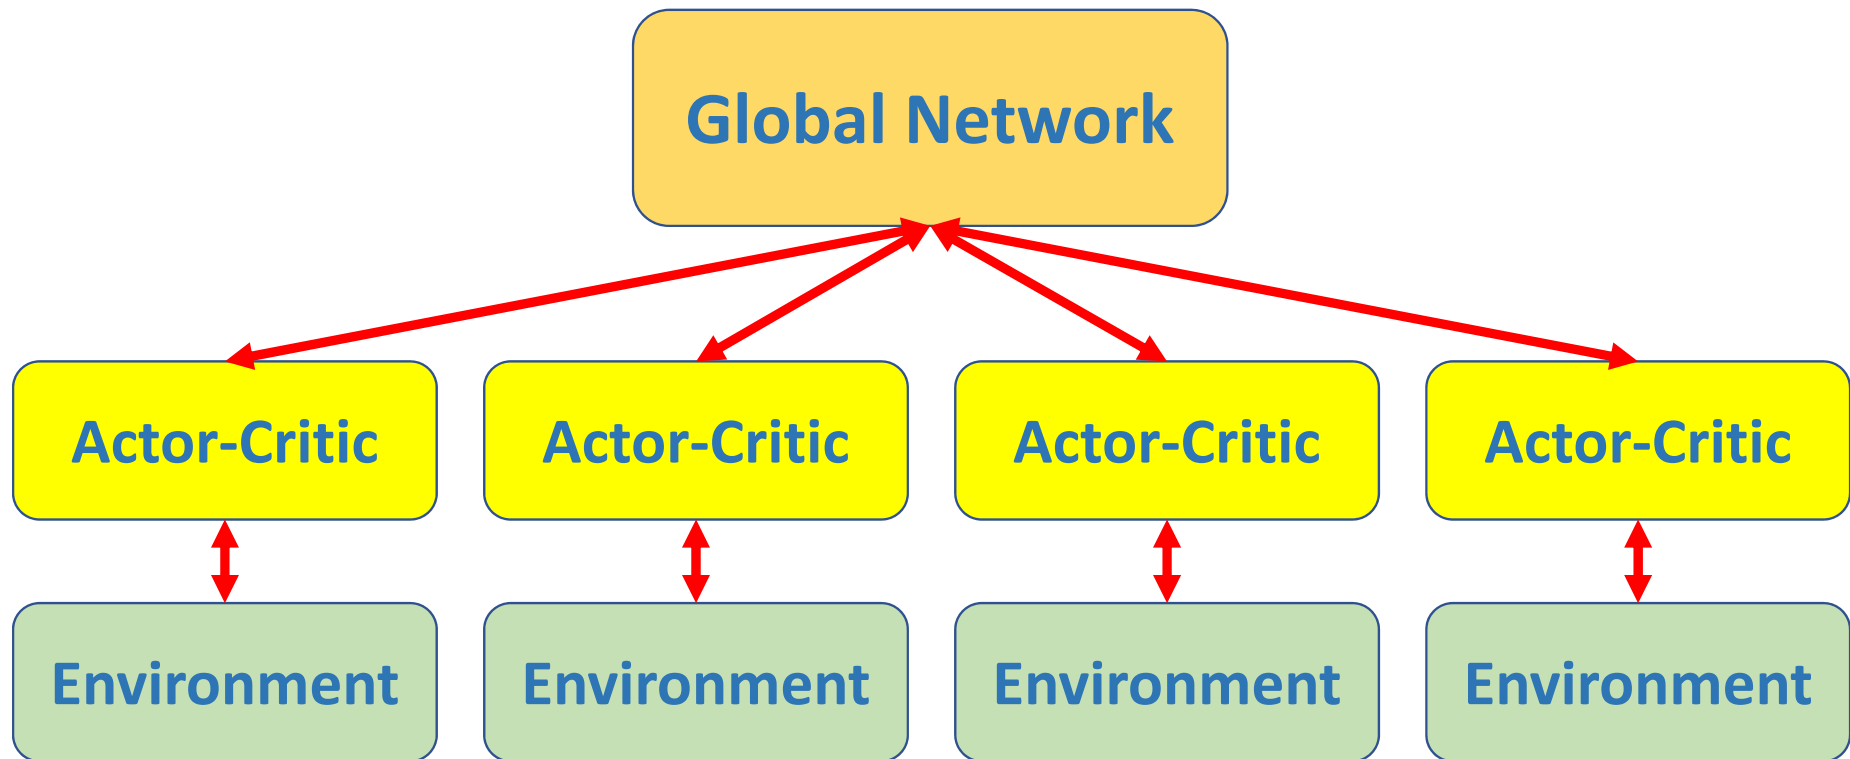




Deep Deterministic Policy Gradient

- ◆ Actor : Policy Gradient
- ◆ Critic : DGN
- ◆ DDPG can only be used for environments with continuous action spaces.

Asynchronous Advantage Actor-Critic (A3C)

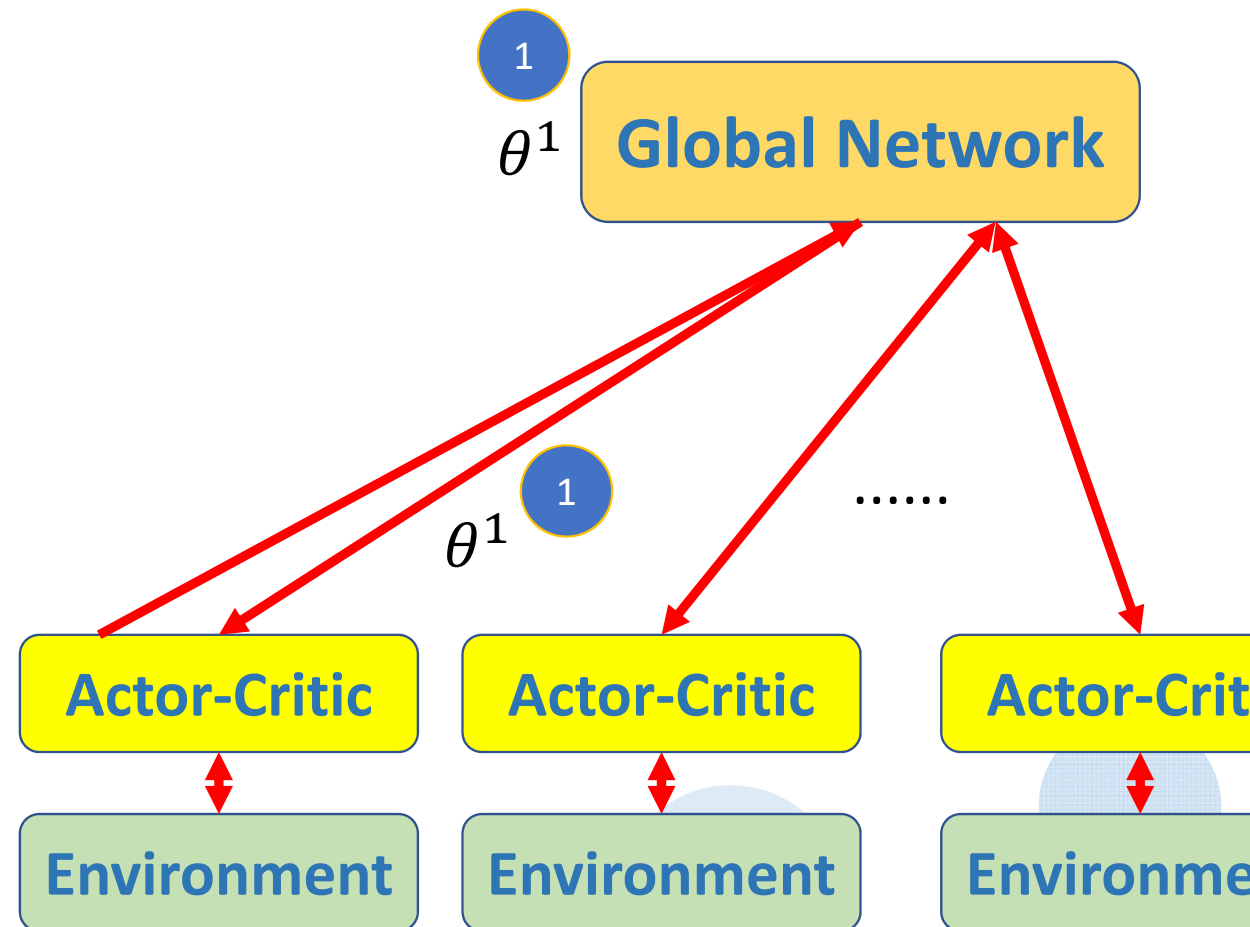


Asynchronous Advantage Actor-Critic (A3C)

- ◆ Train with a lot of Actor-Critic and use a central control for all Actor-Critic.
- ◆ Each Actor-Critic can upload their learned experience to the Global Network, and can also update their parameters using Global Network.

Asynchronous Advantage Actor-Critic (A3C)

- ◆ 1. Copy global parameters
- ◆ 2. Sampling some data
- ◆ 3. Compute gradients
- ◆ 4. Update global models

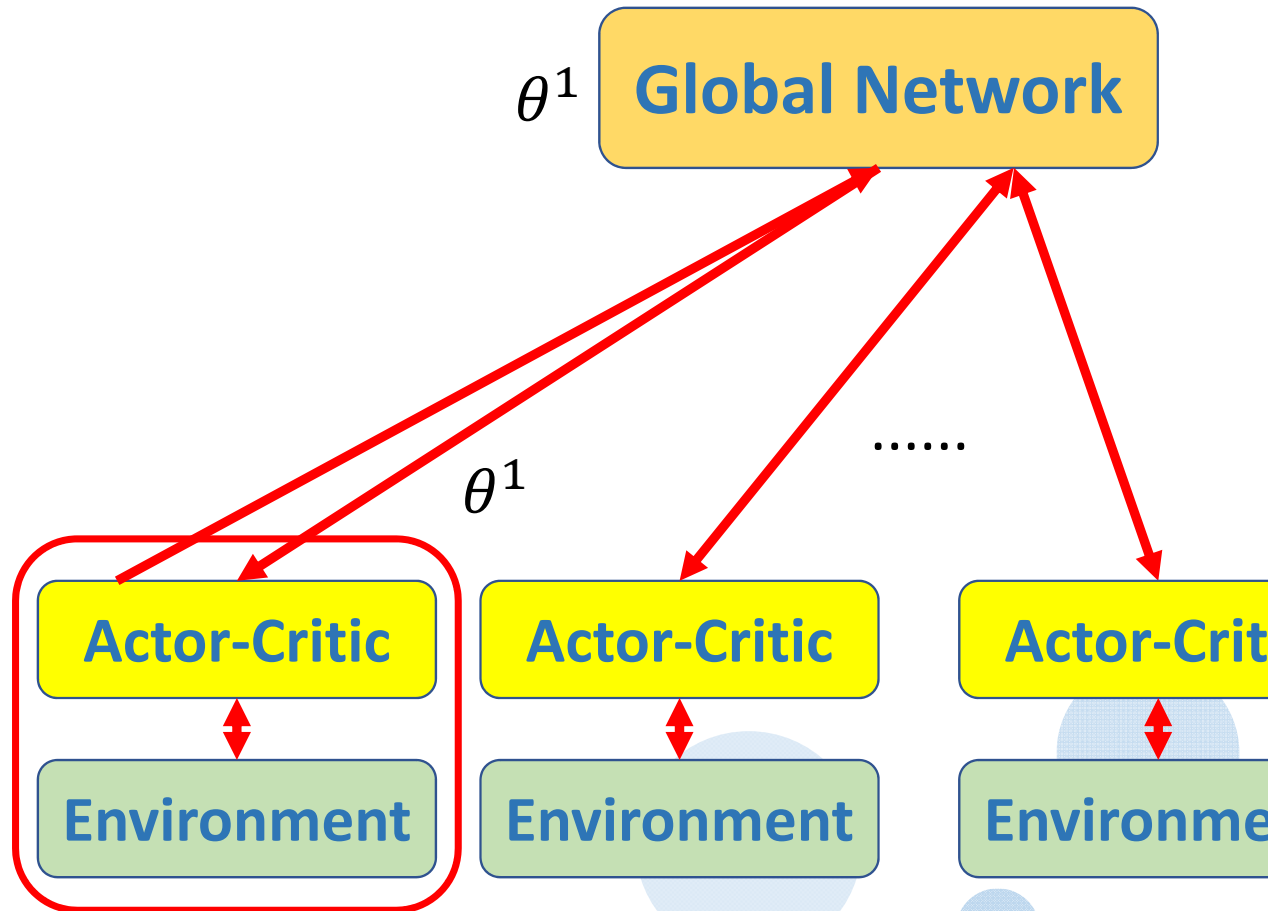


Asynchronous Advantage Actor-Critic (A3C)

- ◆ 1. Copy global parameters
- ◆ 2. **Sampling some data**
- ◆ 3. Compute gradients
- ◆ 4. Update global models

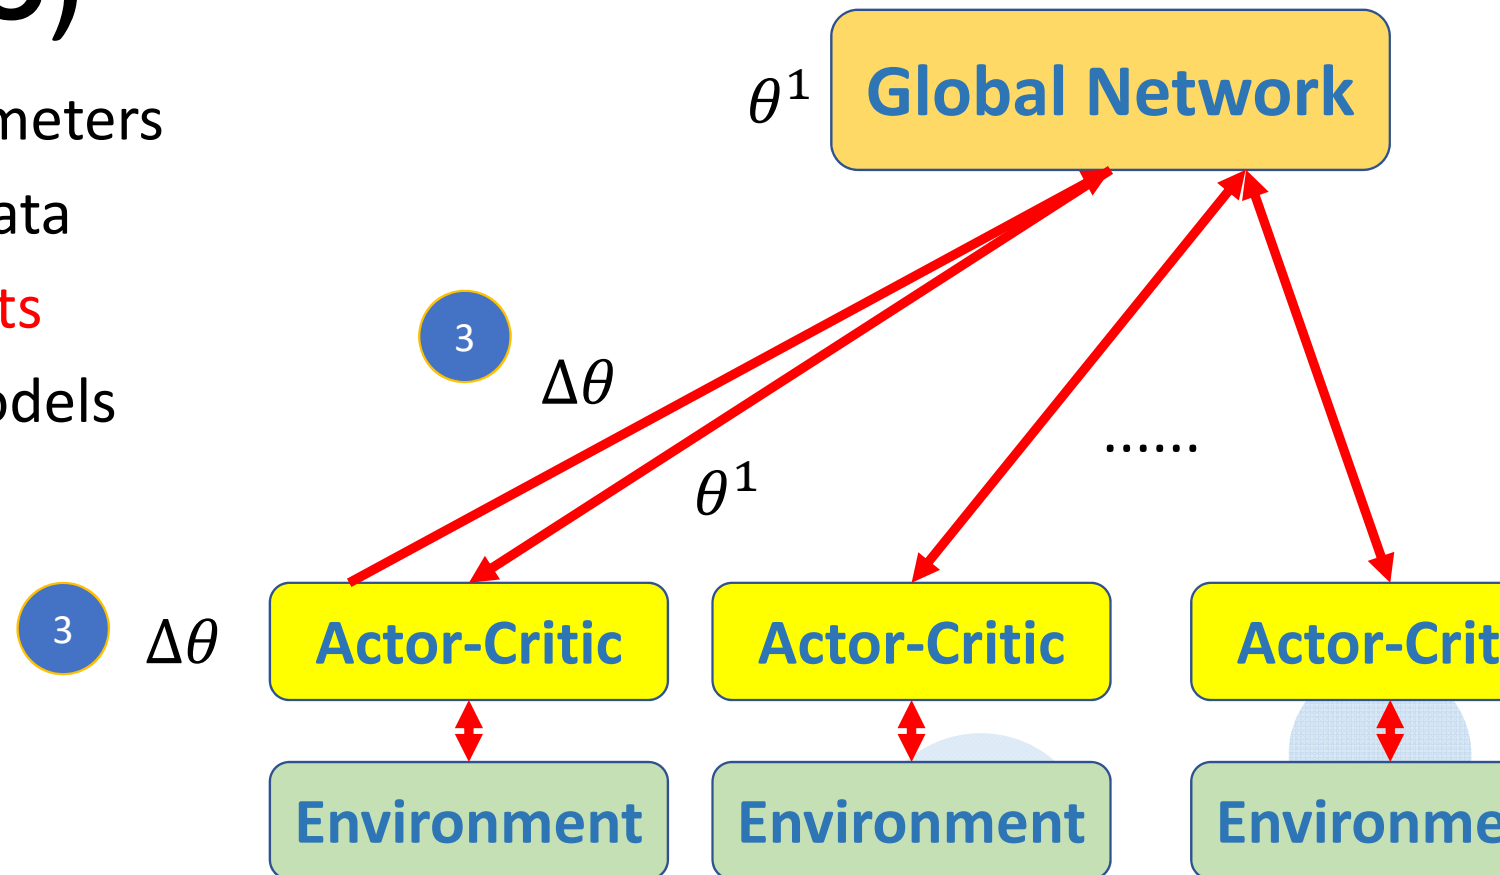
Interact with the environment to obtain information

2



Asynchronous Advantage Actor-Critic (A3C)

- ◆ 1. Copy global parameters
- ◆ 2. Sampling some data
- ◆ 3. **Compute gradients**
- ◆ 4. Update global models

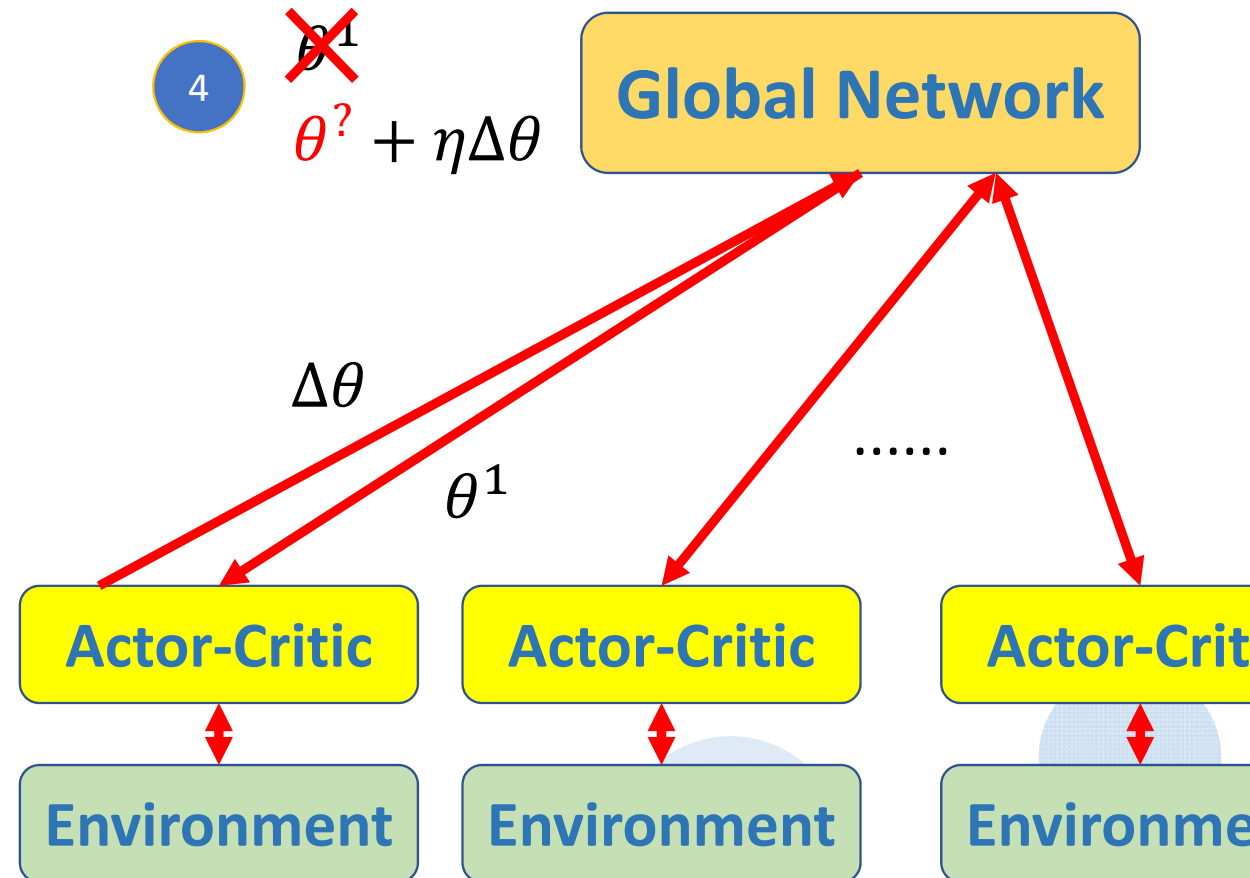


Asynchronous Advantage Actor-Critic (A3C)

- ◆ 1. Copy global parameters
- ◆ 2. Sampling some data
- ◆ 3. Compute gradients
- ◆ 4. **Update global models**


Others may update the Global network, so θ may not be the same

$\Delta\theta$



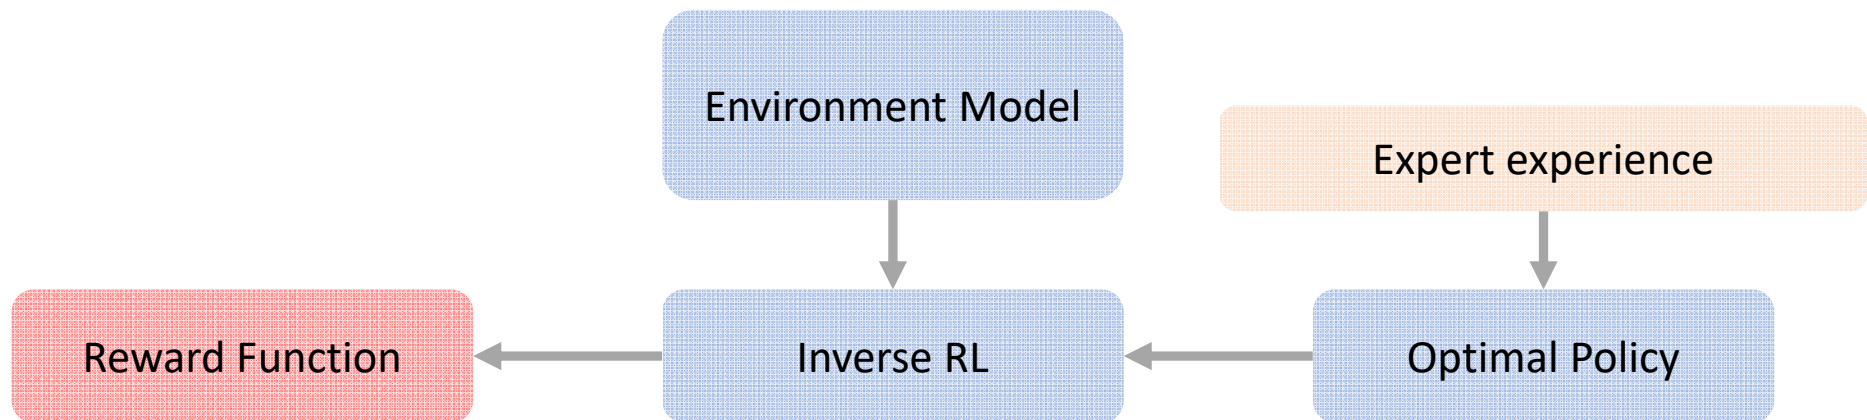
Imitation Learning

- We can learn by imitating the actions of experts.



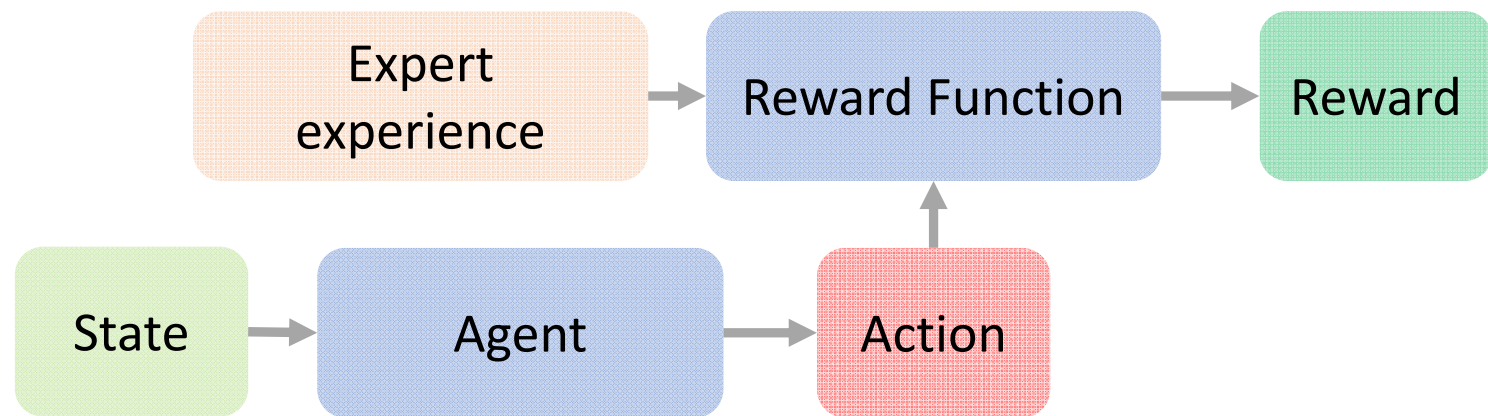
Inverse RL

- We have some expert episodes available.
- Try to find a reward function, which can give expert episode large reward, while our episode generated by our agent low reward.

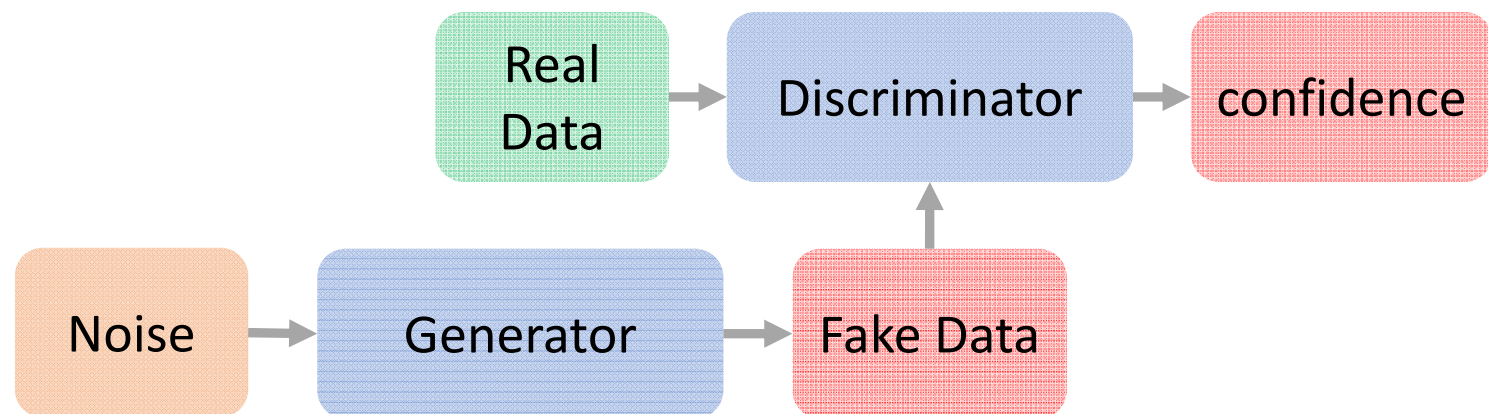


Inverse RL & GAN

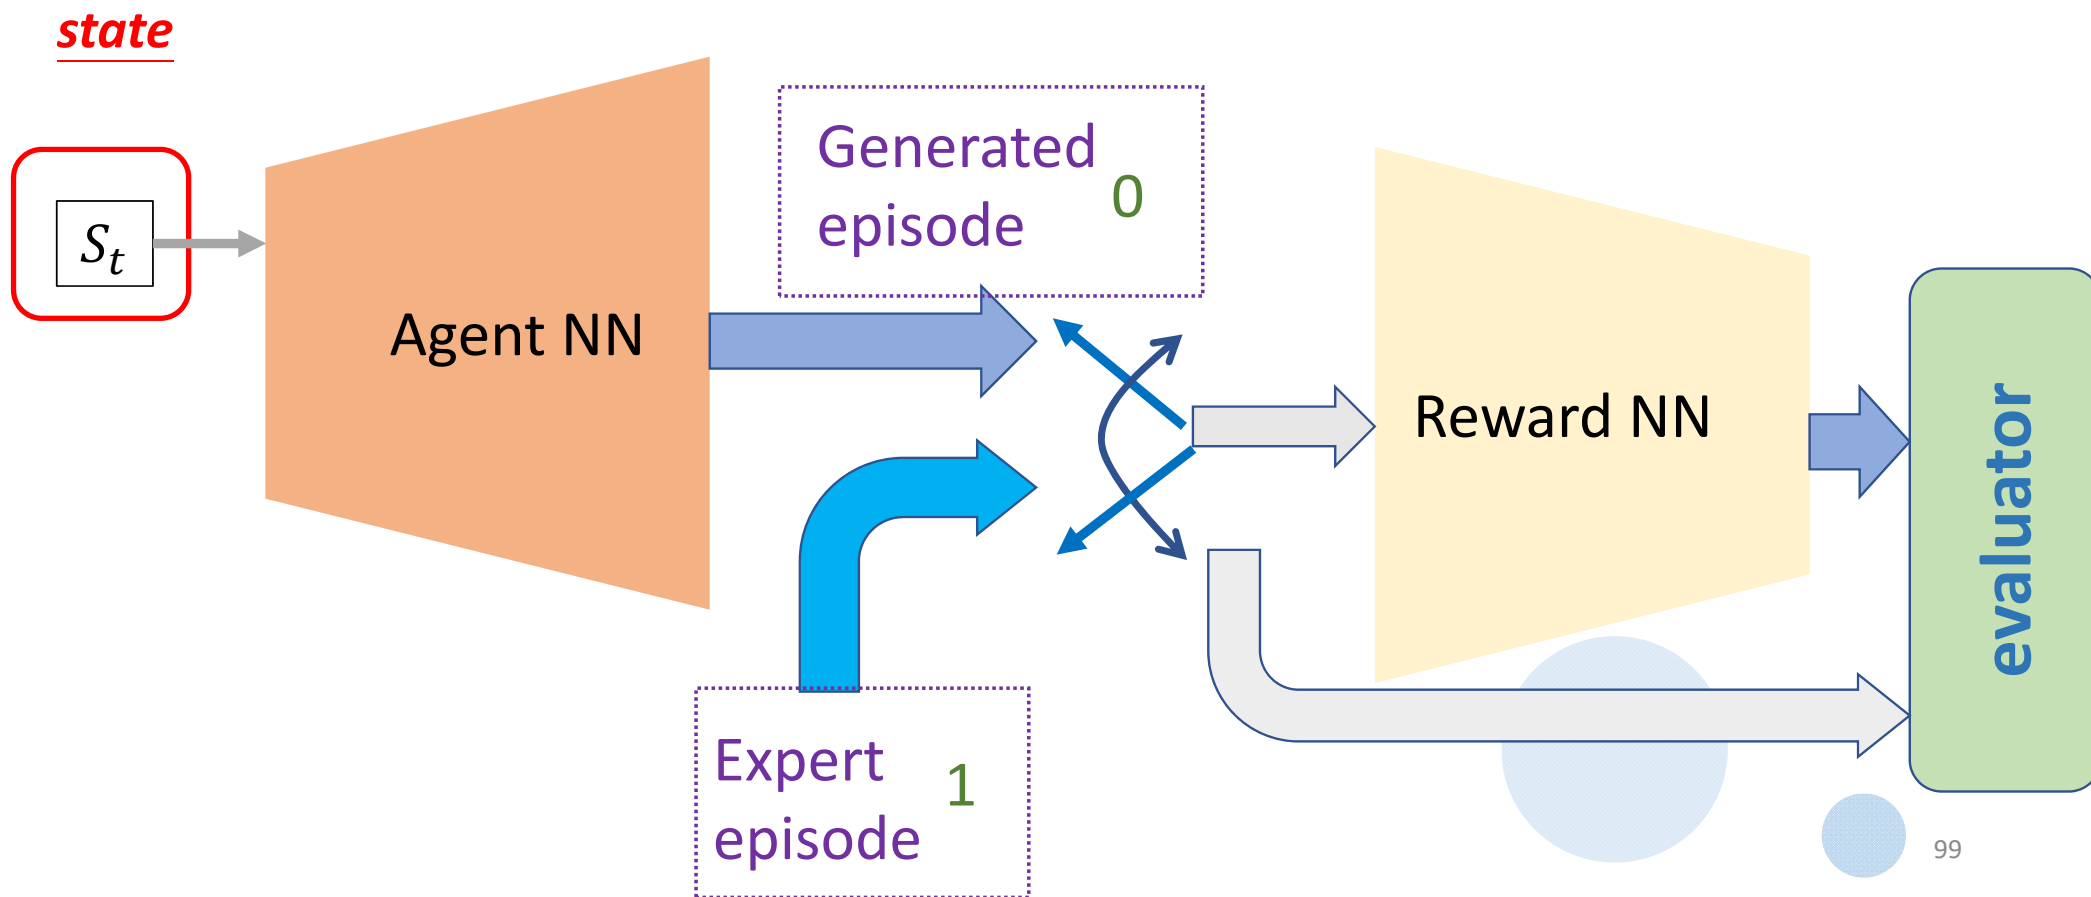
IRL



GAN



Inverse RL





Guided Cost Learning

Guided Cost Learning:
Deep Inverse Optimal Control via Policy Optimization

Chelsea Finn, Sergey Levine, Pieter Abbeel
UC Berkeley

Source <http://rll.berkeley.edu/gcl/>

Thanks for your listening!

