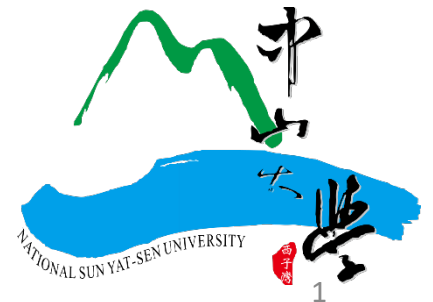


# Image Classification and Nearest Neighbor

Chia-Po Wei

Department of Electrical Engineering  
National Sun Yat-sen University





# Outline

- Image Classification
  - Challenges
  - Machine Learning vs. Classical Programming
- k-Nearest Neighbors (k-NN)
  - Hyperparameters
  - Cross-Validation


# Tasks in Computer Vision

Classification



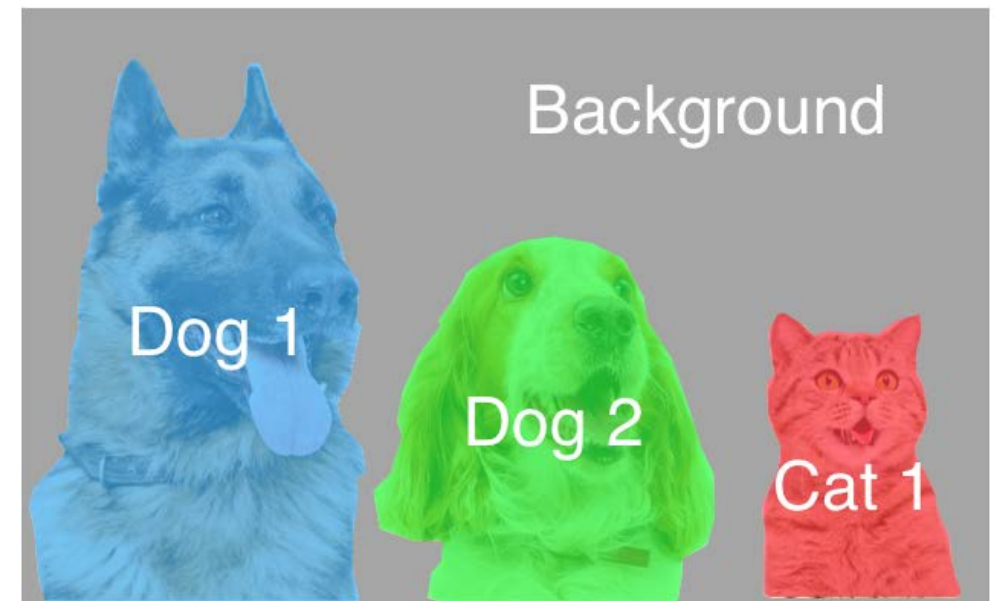
CAT

Object Detection



DOG, DOG, CAT


Instance Segmentation



DOG, DOG, CAT

# Image Classification

- The task of **image classification** is to predict the class label of a test image.
- Example:




Test image



prob.	label
0.7	Lincoln
0.1	Washington
0.1	Jefferson
0.1	Obama

# Images as Matrices





179	167	150	165	156	145	145	183	184
175	172	157	166	169	180	168	168	165
126	164	170	150	121	135	127	112	132
151	130	137	188	196	186	170	178	165
170	166	161	170	155	164	167	173	187
150	154	155	157	161	173	147	164	165
170	149	135	141	151	173	157	170	170
165	174	186	186	166	176	166	168	173
149	165	174	169	152	188	200	191	196
191	175	155	164	164	189	181	153	161

- An image is just a matrix of pixels
- E.g. The left image is 480x600
- The value of each pixel in an image is between 0~255




# Color Images

```
img = cv2.imread('birth_of_venus.jpg')  
print(img.shape) #(480,600,3)
```




- A color image consists of three RGB channels
- The sequence of the channels can be RGB or BGR, depending on the library you use for reading the image.

# Challenges: Viewpoint




# Challenges: Illumination





# Challenges: Occlusion



# Challenges: Deformation





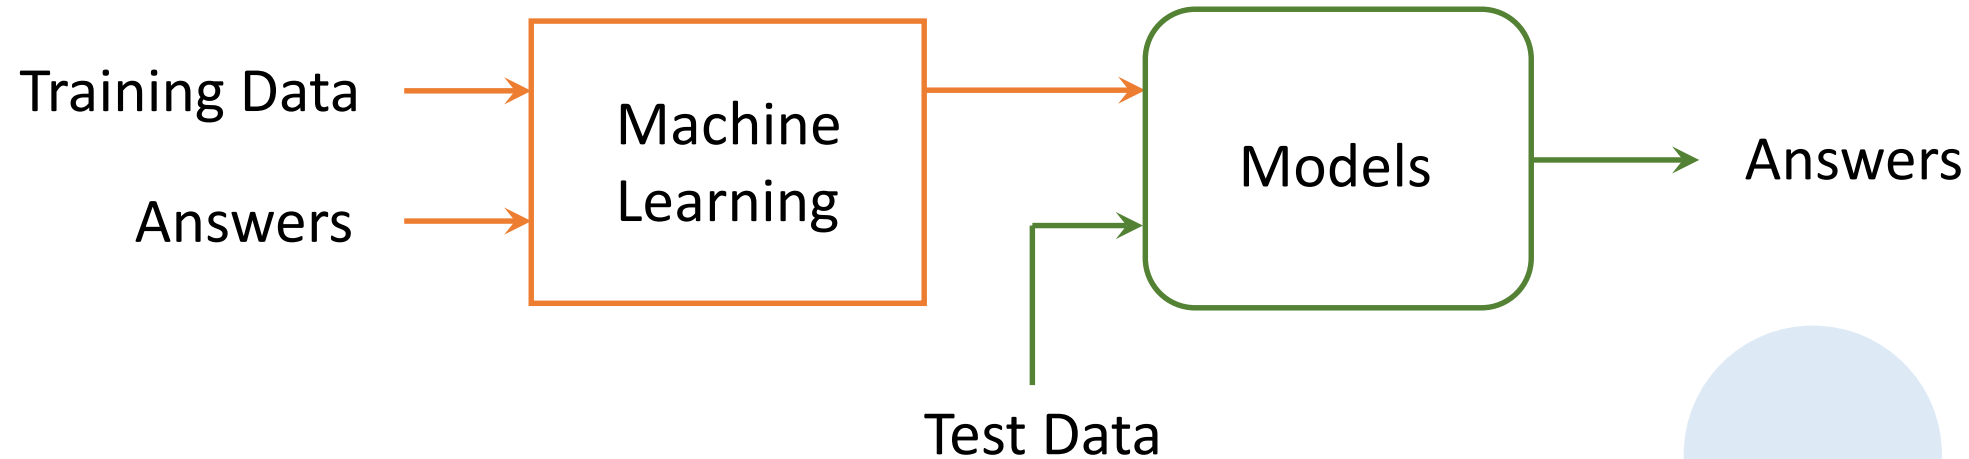

# Challenges: Intra-class Variation



# Challenges: Background Clutter



# Machine Learning vs. Classical Programming





# Machine Learning: Data-Driven Approach to Classification

1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on test images

```
def train(images, labels):  
    # Use machine learning to train a model for classification  
    return classification_model
```


```
def predict(classification_model, test_images):  
    # Use the classification model to predict  
    # the label of each test image  
    return test_labels
```

# Nearest Neighbor (NN)

- Given a set of training images with class labels
- For a test image
  1. Calculate the distance to each training image, and store the distance in a list
  2. Sort the list of distances in ascending order
  3. Assign the class label of the first point in the list to the test image




# k-Nearest Neighbors (k-NN)




- k-NN takes **majority vote** from k closest neighbors instead of copying label from the nearest neighbor.
- In practice, k-NN performs better than NN.
- But how to determine the value of k?

# k-NN Illustration



- The test sample ● should be classified either to Class 1 ■ or to Class 2 ▲
- If  $k = 1$  (inside circle) it is assigned to Class 1 ■ because there is only 1 square.
- If  $k = 3$  (outside circle) it is assigned to Class 2 ▲ because there are 2 triangles and only 1 square.

# Quiz 1




1. What is the label of the green target if using 4-nn as the classifier?
2. What can we do if the class votes are tied?




# Choices of $k$


$k = 1$



$k = 3$



$k = 5$




- For NN classifiers, outlier points create small islands of likely incorrect predictions
- 5-NN classifiers smooth over these irregularities, likely leading to better **generalization** on the test data (not shown).

# Common Distance Metric

L1 (Manhattan) distance


$$\text{dist}(\mathbf{I}_1, \mathbf{I}_2) = \sum_k |\mathbf{I}_1^k - \mathbf{I}_2^k|$$



$$\|\mathbf{x}\|_1 = |x_1| + |x_2| = 1$$

L2 (Euclidean) distance

$$\text{dist}(\mathbf{I}_1, \mathbf{I}_2) = \sqrt{\sum_k (\mathbf{I}_1^k - \mathbf{I}_2^k)^2}$$



$$\|\mathbf{x}\|_2 = \sqrt{x_1^2 + x_2^2} = 1$$

# L1 Distance Metric

$$\text{dist}(\mathbf{I}_1, \mathbf{I}_2) = \sum_k |\mathbf{I}_1^k - \mathbf{I}_2^k|$$

Test image

23	155	166	45
34	203	200	63
66	255	195	98
77	143	150	88

Training image

45	90	78	42
87	122	144	94
76	156	167	83
34	234	255	72

—

=


Absolute difference

22	65	88	3
53	81	56	31
10	99	28	15
43	91	105	16

sum  
→ 806

# Minkowski Distance

$$\text{dist}(\mathbf{I}_1, \mathbf{I}_2) = \left( \sum_k (\mathbf{I}_1^k - \mathbf{I}_2^k)^p \right)^{1/p}$$



# Hyperparameters

- For k-NN
  - What is the best value of k to use?
  - What is the best distance metric to use? (L1, L2, or dot products)
- These choices are called **hyperparameters**.
  - Come up in the design of Machine Learning algorithms
  - Choices regarding the algorithm that we set rather than learn



# Choosing Hyperparameters

**Method 1:** Split data into **train** and **test**. Choose hyperparameters that work best on **test**.



The problem of Method 1 is that we have no idea how the algorithm will perform on new data.

**Method 2:** Split data into **train**, **val**, and **test**. Choose hyperparameters on **val** and evaluate on **test**.



# Cross-Validation for Hyperparameter Tuning

**Method 3 (Cross-Validation):** Split training data into folds. Try one fold as **val** and the remaining folds as **train**. Average the results of each fold as **val** to choose hyperparameters.


- Below is an example of 3-fold cross-validation. Blue folds are for training, while green folds are for validation.

fold 1	fold 2	fold 3	test
--------	--------	--------	------

fold 1	fold 2	fold 3	test
--------	--------	--------	------

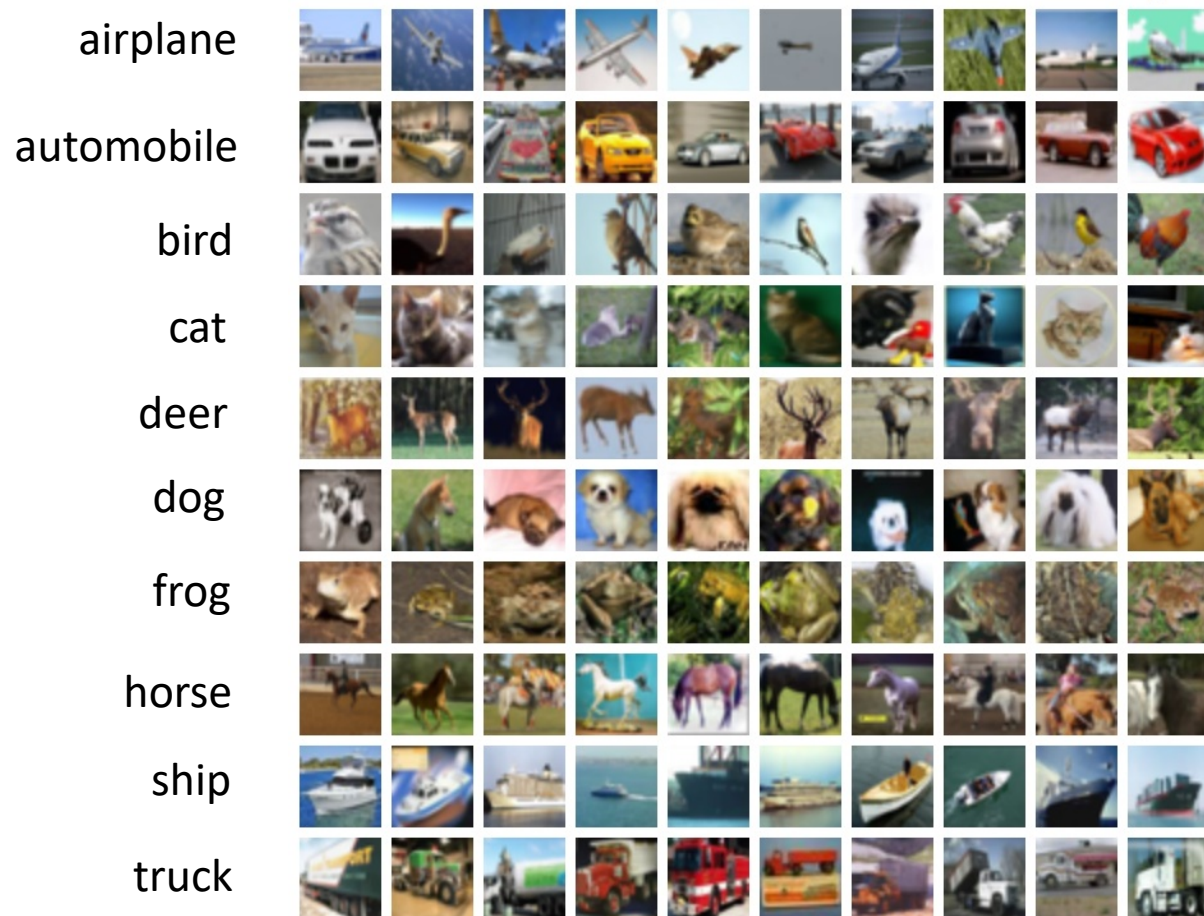
fold 1	fold 2	fold 3	test
--------	--------	--------	------

# Results of 5-fold Cross-Validation



- Each point is a single outcome.
- Each k has 5 points because of 5-fold cross-validation.
- The line goes through the mean, while bars indicated the standard deviation.

# CIFAR-10 Dataset



- 60,000 color images in 10 classes
- Each class has 6,000 images
- Each color image is 32x32
- 50,000 images for training
- 10,000 images for testing
- No overlap between automobiles and trucks

# Class Labels

Original (k-NN)

class	label
airplane	0
automobile	1
bird	2
cat	3
deer	4
dog	5
frog	6
horse	7
ship	8
truck	9

One-Hot Encoding (softmax)

class	label
airplane	[1 0 0 0 0 0 0 0 0 0]
automobile	[0 1 0 0 0 0 0 0 0 0]
bird	[0 0 1 0 0 0 0 0 0 0]
cat	[0 0 0 1 0 0 0 0 0 0]
deer	[0 0 0 0 1 0 0 0 0 0]
dog	[0 0 0 0 0 1 0 0 0 0]
frog	[0 0 0 0 0 0 1 0 0 0]
horse	[0 0 0 0 0 0 0 1 0 0]
ship	[0 0 0 0 0 0 0 0 1 0]
truck	[0 0 0 0 0 0 0 0 0 1]



# Classification Examples of kNN



- Red boxes indicate incorrect results, while blue boxes indicate correct results.
- The number above each image represents the class label of the image.

# Quiz 2

1. What is the class label of the following test image if using NN as the classifier?

Test image



Frontal 1 (f1)

Label 1



Profile 1 (p1)


Label 2



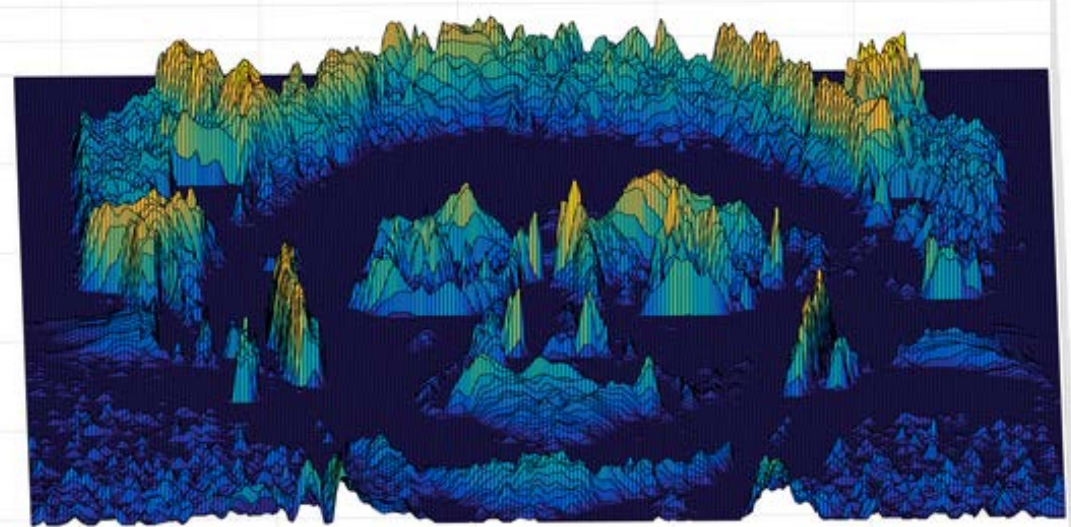
Frontal 2 (f2)

# Quiz 2 (cont.)

$\text{abs}(f1 - p1)$



$\text{abs}(f1 - f2)$



- $\text{sum}(\text{abs}(f1-p1)) / N = 57.7$
- $\text{sum}(\text{abs}(f1-f2)) / N = 11.1$
- $N$  is the number of pixels



# Pros and Cons of kNN

- Pros
  - Simple to implement and understand
  - No need to train a model
- Cons
  - Need to store all the training data
  - Heavy computational cost at test time as the number of training images increases
  - Using pixel differences to compare images is inadequate