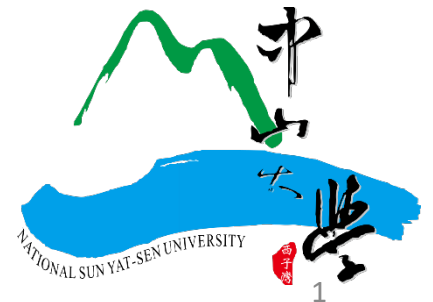


TensorFlow Tutorial

Chia-Po Wei

Department of Electrical Engineering
National Sun Yat-sen University





Outline

- Deep Learning Framework
- TensorFlow Core Walkthrough
 - Tensor, Graph, Session
 - `tf.data.Dataset`
 - Symbolic Execution vs. Eager Execution
- Fitting a Linear Model

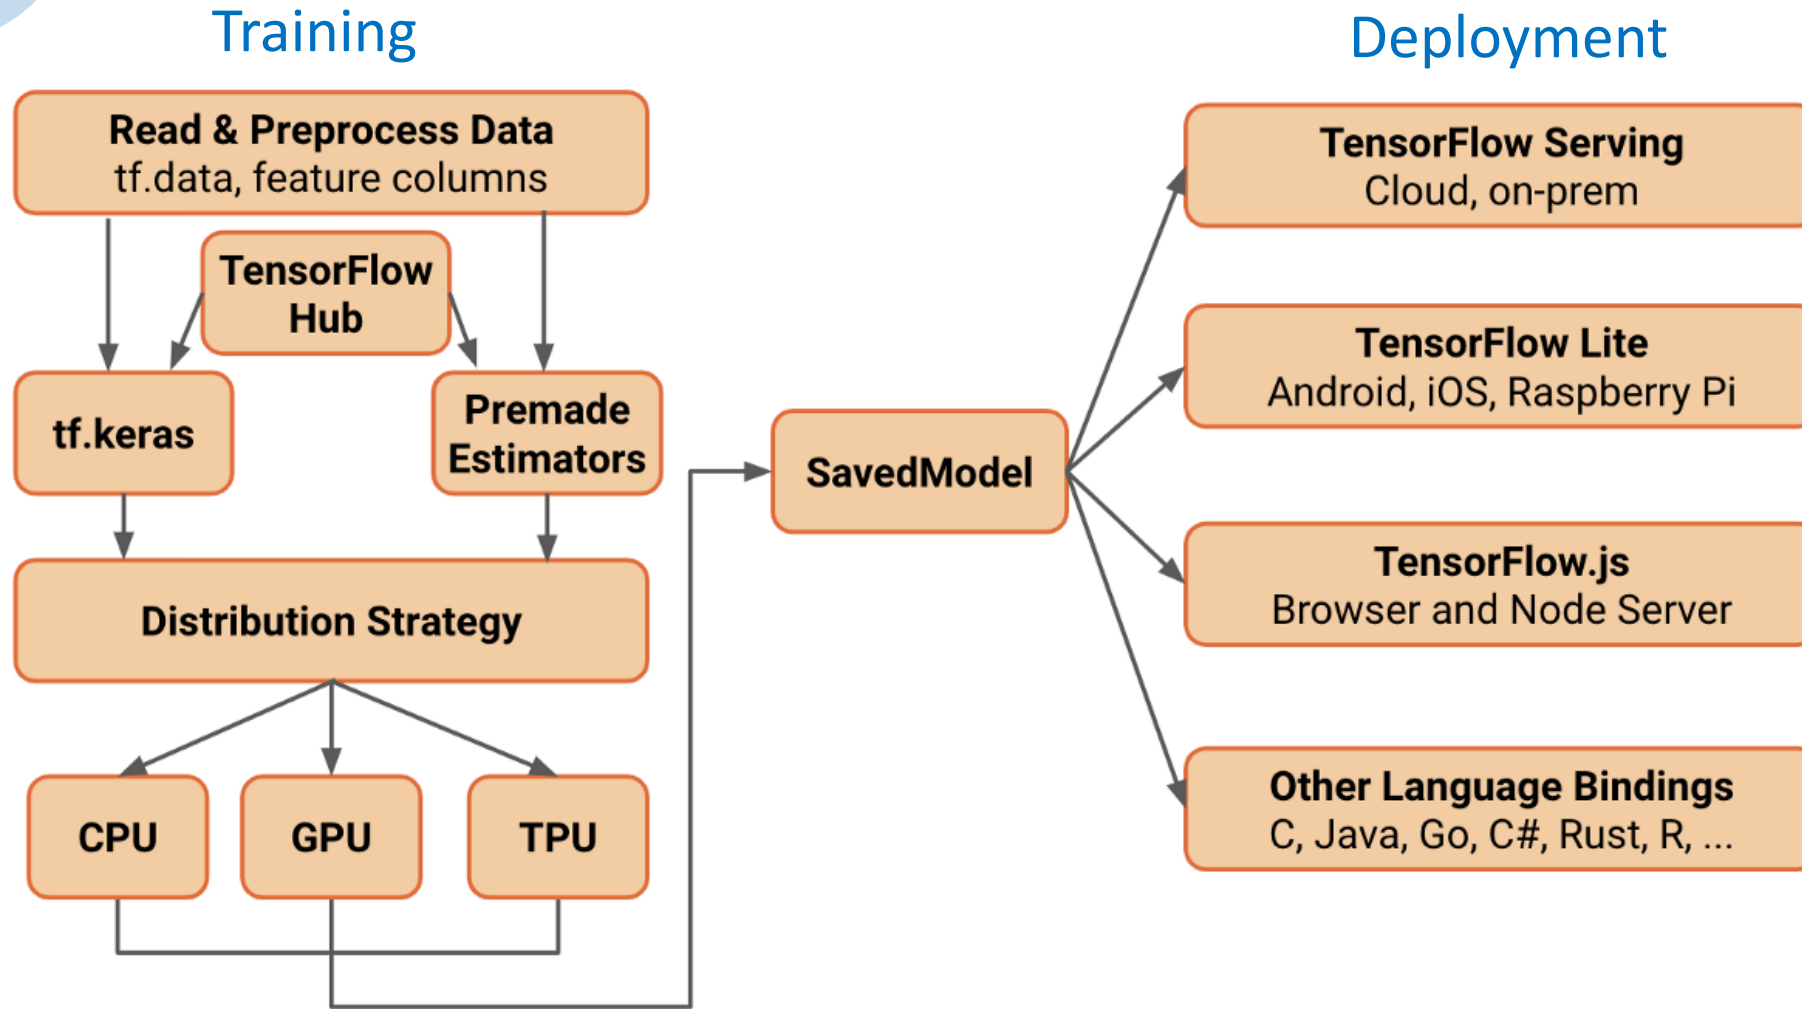
What is TensorFlow?

- TensorFlow is an end-to-end open source machine learning platform.
- Tensorflow helps you develop and train Machine Learning models.



TensorFlow

Tensorflow 2.0 (Conceptual Diagram)




Deep Learning Frameworks



- Tensorflow (Google), Pytorch (Facebook), Caffe2 (Facebook), mxnet (Apache), CNTK (Microsoft)

ONNX (Open Neural Network Exchange)

- ONNX is a open format to represent deep learning models.
- ONNX allows you to switch between deep learning frameworks.
- Facebook, Microsoft, Amazon announced support for ONNX, but Google hasn't.




ONNX

TensorFlow vs. Numpy

- Both TensorFlow and Numpy provide support for large multi-dimensional (N-d) arrays with a collection of mathematical functions.
- But Numpy cannot automatically compute derivatives and cannot take advantage of GPU.



vs.



Numpy and TensorFlow Comparison

Numpy	TensorFlow
<code>a = np.zeros((2,2)); b = np.ones((2,2))</code>	<code>a = tf.zeros((2,2)); b = tf.ones((2,2))</code>
<code>np.sum(b, axis=1)</code>	<code>tf.reduce_sum(b, axis=1)</code>
<code>a.shape</code>	<code>a.get_shape()</code> or <code>tf.shape(a)</code>
<code>np.reshape(a, (1,5))</code>	<code>tf.reshape(a, (1,5))</code>
<code>a * 7 + 1</code>	<code>a * 7 + 1</code>
<code>np.dot(a, b)</code>	<code>tf.matmul(a, b)</code>
<code>a[0, 1], a[:, 1], a[0,:]</code>	<code>a[0, 1], a[:, 1], a[0,:]</code>

How to install TensorFlow?

- After installing **anaconda (Python 3.7)**
 - CPU version: `conda install tensorflow`
 - GPU version: `conda install tensorflow-gpu`
 - Prefer `conda install` to `pip install` if you want to install the GPU version.
 - Because you don't have to struggle to find the compatible NVIDIA GPU drivers and CUDA Toolkits. `conda` will do the job for you.
- It is recommended that you learn to setup a virtual environment using `conda`.
- This allows you to switch between DL frameworks easily.



Setup your Python environment

```
from __future__ import absolute_import  
from __future__ import division  
from __future__ import print_function
```

```
import numpy as np  
import tensorflow as tf
```

```
print(tf.__version__) # print the version of TensorFlow
```

Tensor: rank, shape

- The central unit of data in TensorFlow is called the tensor.
- A tensor's **rank** is its number of dimensions.
- A tensor's **shape** is a tuple of integers specifying the array's length along each dimension.

3. # rank 0 tensor; a scalar with shape []

[1., 2., 3.] # a rank 1 tensor; a vector with shape [3]

[[1., 2., 3.], [4., 5., 6.]] # a rank 2 tensor; a matrix with shape [2, 3]

[[[1., 2., 3.], [2., 8., 9.]]] # a rank 3 tensor with shape [2, 1, 3]

Quiz 1

Determine the rank and shape of the following Tensors.

1. `[[[1, 2, 3], [4, 5, 6]]]`
2. `[[[1], [3], [5]], [[2], [4], [6]]]`

Tensor Examples

- A tensor has a **data type** (float32, int32, string, etc.) and a **shape**.
- Each element in a tensor has **the same** data type, and the data type is known.

Four types of tensors:

- `tf.Variable`
- `tf.constant`
- `tf.placeholder`
- `tf.SparseTensor`

Rank 0 examples:

- `mammal = tf.Variable("Elephant", tf.string)`
- `ignition = tf.Variable(451, tf.int16)`
- `floating = tf.Variable(3.14159265359, tf.float64)`
- `complicated = tf.Variable(12.3 - 4.85j, tf.complex64)`

Rank 1 examples:

- `mammal = tf.Variable(["Elephant"], tf.string)`
- `ignition = tf.Variable([123, 456], tf.int16)`

TensorFlow Core Walkthrough

- A typical TensorFlow program consists of two discrete sections:
 - Building the computational graph (`tf.Graph`)
 - Running the computational graph (using `tf.Session`)
- A graph is composed of
 - `tf.Operation` (the nodes): Operations that consume and produce tensors.
 - `tf.Tensor` (the edges): Represent the values that will flow through the graph.

★ `tf.Tensors` do not have values, they are just handles to elements in the graph.

Simple Graph

```
a = tf.constant(3.0, dtype=tf.float32)
b = tf.constant(4.0) # also tf.float32 implicitly
total = a + b
print(a)
print(b)
print(total)
```

Output:

```
Tensor("Const:0", shape=(), dtype=float32)
Tensor("Const_1:0", shape=(), dtype=float32)
Tensor("add:0", shape=(), dtype=float32)
```



Session

- To evaluate tensors, a `tf.Session` object is instantiated.
- If a `tf.Graph` is like a .py file, a `tf.Session` is like the python executable.

```
a = tf.constant(3.0, dtype=tf.float32)
b = tf.constant(4.0)
total = a + b
sess = tf.Session()
print(sess.run(total))
print(sess.run({'ab': (a, b), 'total': total}))
```

Output:

7.0

{'total': 7.0, 'ab': (3.0, 4.0)}

Placeholder for Feeding Data

- A graph can be parameterized to accept external inputs, known as **placeholders**.
- A **placeholder** is a promise to provide a value later.

```
x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = x + y
sess = tf.Session()
print(sess.run(z, feed_dict={x: 3, y: 4.5}))
print(sess.run(z, feed_dict={x: [1, 3], y: [2, 4]}))
```

Output:

7.5
[3. 7.]

tf.data.Dataset

- Placeholders work for simple experiments, but `tf.data` are the **preferred** method of streaming data into a model.
- To get a runnable `tf.Tensor` from a Dataset you must first convert it to a `tf.data.Iterator`, and then call the Iterator's `tf.data.Iterator.get_next` method.

```
my_data = [[0, 1,], [2, 3,], [4, 5,], [6, 7,]]  
slices = tf.data.Dataset.from_tensor_slices(my_data)  
next_item =  
slices.make_one_shot_iterator().get_next()
```

```
sess = tf.Session()  
while True:  
    try:  
        print(sess.run(next_item))  
    except tf.errors.OutOfRangeError:  
        break
```

tf.data.Dataset (cont.)

Read records from a list of files.

```
dataset = TFRecordDataset(["file1.tfrecord", "file2.tfrecord", ...])
```

Parse string values into tensors.

```
dataset = dataset.map(lambda record: tf.parse_single_example(record, ...))
```

Randomly shuffle using a buffer of 10000 examples.

```
dataset = dataset.shuffle(10000)
```

Repeat for 100 epochs.

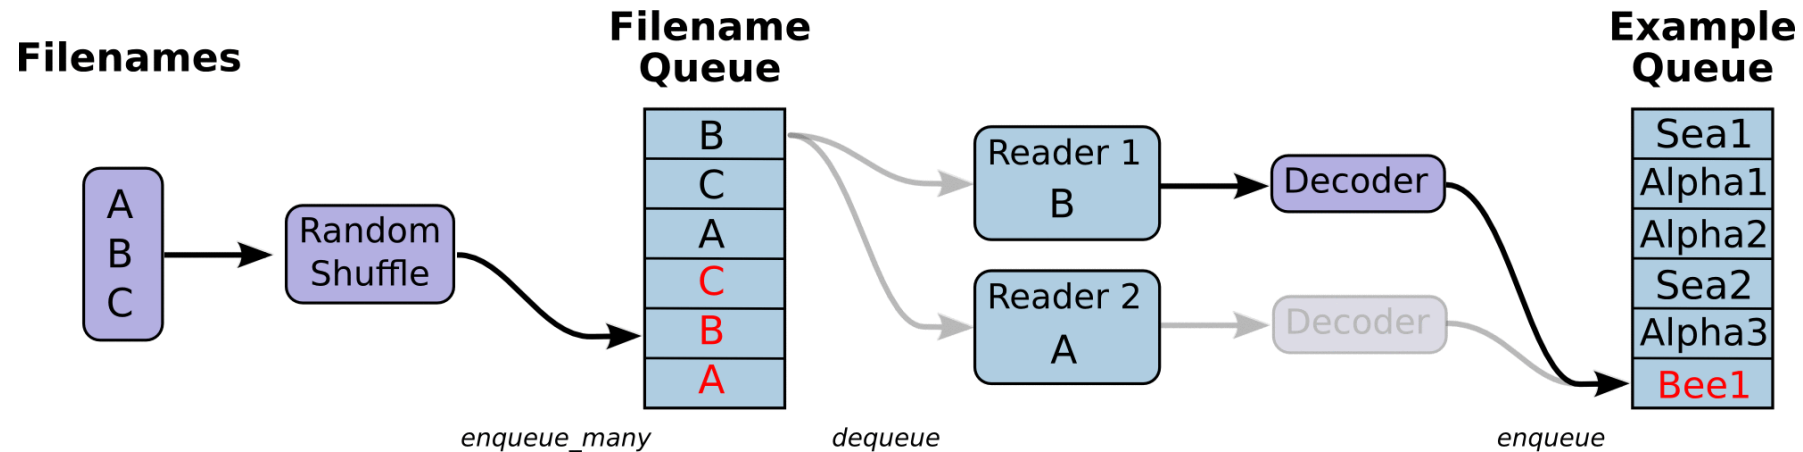
```
dataset = dataset.repeat(100)
```

Combine 128 consecutive elements into a batch.

```
dataset = dataset.batch(128)
```

- `tf.data.Dataset` has been introduced since the release of TensorFlow 1.4.

Queue-based Pipelines (DEPRECATED)



- `tf.train.string_input_producer` outputs strings to a queue for an input pipeline.
- Queue-based input pipelines have been replaced by `tf.data`.

How are Machine Learning Models represented?

Model is a **Data Structure**
(e.g. a Graph)

aka

- Symbolic
- Deferred Execution
- Defined-and-run

Model is a **Program**
(e.g. Python Code)

aka

- Imperative
- Eager Execution
- Defined-by-run


Why Symbolic Execution?

Pros (Makes (de)serialization easier)

- Parallelism
- Distributed execution (CPUs, GPUs, TPUs)
- Compilation (XLA compiler)
- Portability (language independent)

Cons (Less intuitive)

- Harder to debug
- Harder to write control flow structures
- Harder to write dynamic models



Why Eager Execution?

- **Eager execution** evaluates operations **immediately, without building graphs**: operations return concrete values instead of constructing a graph to run later.
- **Advantages**
 - **An intuitive interface**—Structure your code naturally and use Python data structures. Quickly iterate on small models and small data.
 - **Easier debugging**—Call ops directly to inspect running models and test changes. Use standard Python debugging tools for immediate error reporting.
 - **Natural control flow**—Use Python control flow instead of graph control flow, simplifying the specification of dynamic models.

Eager Execution (cont.)

```
a = tf.constant(3.0, dtype=tf.float32)
b = tf.constant(4.0)
total = a + b
print(total)
```

Output:

```
Tensor("add:0", shape=(), dtype=float32)
```

```
tf.enable_eager_execution()
```


```
a = tf.constant(3.0, dtype=tf.float32)
b = tf.constant(4.0)
total = a + b
print(total)
```

Output:

```
tf.Tensor(7.0, shape=(), dtype=float32)
```

```
# try print(total.numpy())
```



Static Model Structures



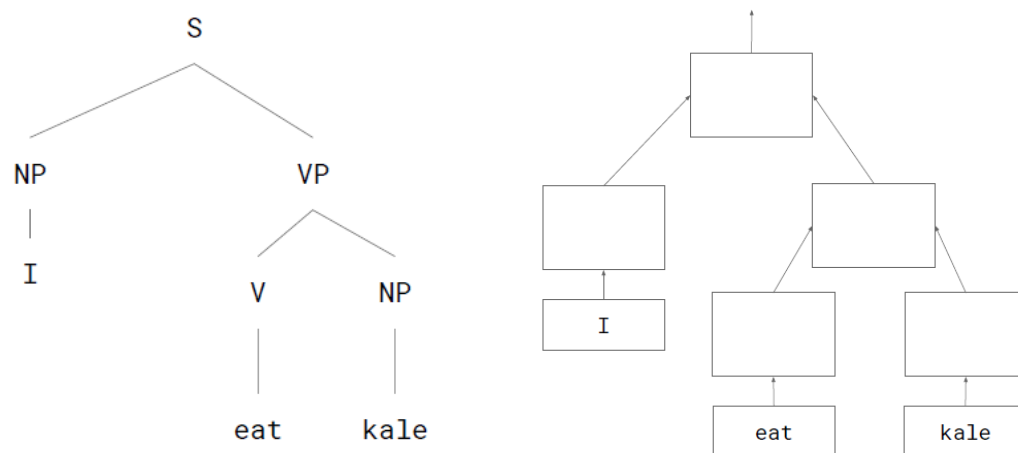
- Model structure is fixed regardless of input data
- The majority of DL models for image, audio and numerical data

Dynamic Model Structures

Traditional RNN



Tree RNN (Dynamic Models)



- Structures that change a lot with input data
- Can deal with hierarchical structures in natural language
- Difficult to write in the symbolic way (using `tf.cond` and `tf.while_loop`)
- Straightforward with Eager (Using the native Python control flow)


Fitting a Linear Model

- We will walk through the example of a simple linear model $f(\mathbf{x}) = \mathbf{W}\mathbf{x} + \mathbf{b}$, which has two variables \mathbf{W} and \mathbf{b} .

1. Obtain training data
2. Define the model
3. Define a loss function
4. Run through the training data

Step 1. Obtaining Training Data

- Synthesize data such that a well trained model would have $W = 3.0$ and $b = 2.0$



`TRUE_W = 3.0`

`TRUE_b = 2.0`

`NUM_EXAMPLES = 1000`

`inputs = tf.random_normal(shape=[NUM_EXAMPLES])`

`noise = tf.random_normal(shape=[NUM_EXAMPLES])`

`outputs = inputs * TRUE_W + TRUE_b + noise`

- Given inputs and outputs, our goal is to learn W and b such that $f(\text{inputs}) = \text{outputs}$.

Step 2. Define the Model

```
class Model(object):  
    def __init__(self):  
        self.W = tf.Variable(-1.0)  
        self.b = tf.Variable(1.0)  
  
    def __call__(self, x):  
        return self.W * x + self.b
```

```
model = Model()  
assert model(3.0).numpy() == 15.0
```



$$f(\mathbf{x}) = \mathbf{W}\mathbf{x} + \mathbf{b}$$

Step 3. Define the Loss Function


```
def loss(predicted_y, desired_y):  
    return tf.reduce_mean(tf.square(predicted_y - desired_y))  
  
def train(model, inputs, outputs, learning_rate):  
    with tf.GradientTape() as tape:  
        current_loss = loss(model(inputs), outputs)  
        dW, db = tape.gradient(current_loss, [model.W, model.b])  
        model.W.assign_sub(learning_rate * dW) # W -= learning_rate * dW  
        model.b.assign_sub(learning_rate * db) # b -= learning_rate * db  
    return current_loss
```

Step 4. Run Through the Training Data

```
model = Model()
Ws, bs = [], []
epochs = range(10)
for epoch in epochs:
    current_loss = train(model, inputs, outputs, learning_rate=0.1)
    Ws.append(model.W.numpy())
    bs.append(model.b.numpy())
    print(epoch, Ws[-1], bs[-1], current_loss)
```

Simulation Results

```
Epoch 0: W=-1.00 b=1.00, loss=18.29831
Epoch 1: W=-0.19 b=1.18, loss=12.15019
Epoch 2: W=0.45 b=1.33, loss=8.18830
Epoch 3: W=0.97 b=1.45, loss=5.63520
Epoch 4: W=1.39 b=1.55, loss=3.98991
Epoch 5: W=1.72 b=1.63, loss=2.92962
Epoch 6: W=1.99 b=1.69, loss=2.24632
Epoch 7: W=2.20 b=1.74, loss=1.80596
Epoch 8: W=2.37 b=1.79, loss=1.52217
Epoch 9: W=2.51 b=1.82, loss=1.33927
      ⋮
Epoch 29: W=3.07 b=1.96, loss=1.00772
```



- Blue dots are training data
- Red dots are prediction results

References

- TensorFlow Introduction (https://www.tensorflow.org/guide/low_level_intro)
- Tensors (<https://www.tensorflow.org/guide/tensors>)
- Graphs and Sessions (<https://www.tensorflow.org/guide/graphs>)
- Eager Execution (<https://www.tensorflow.org/guide/eager>)
- Fitting a Linear Model (https://www.tensorflow.org/tutorials/eager/custom_training)