


# Introduction to Neural Networks

Chia-Po Wei

Department of Electrical Engineering  
National Sun Yat-sen University

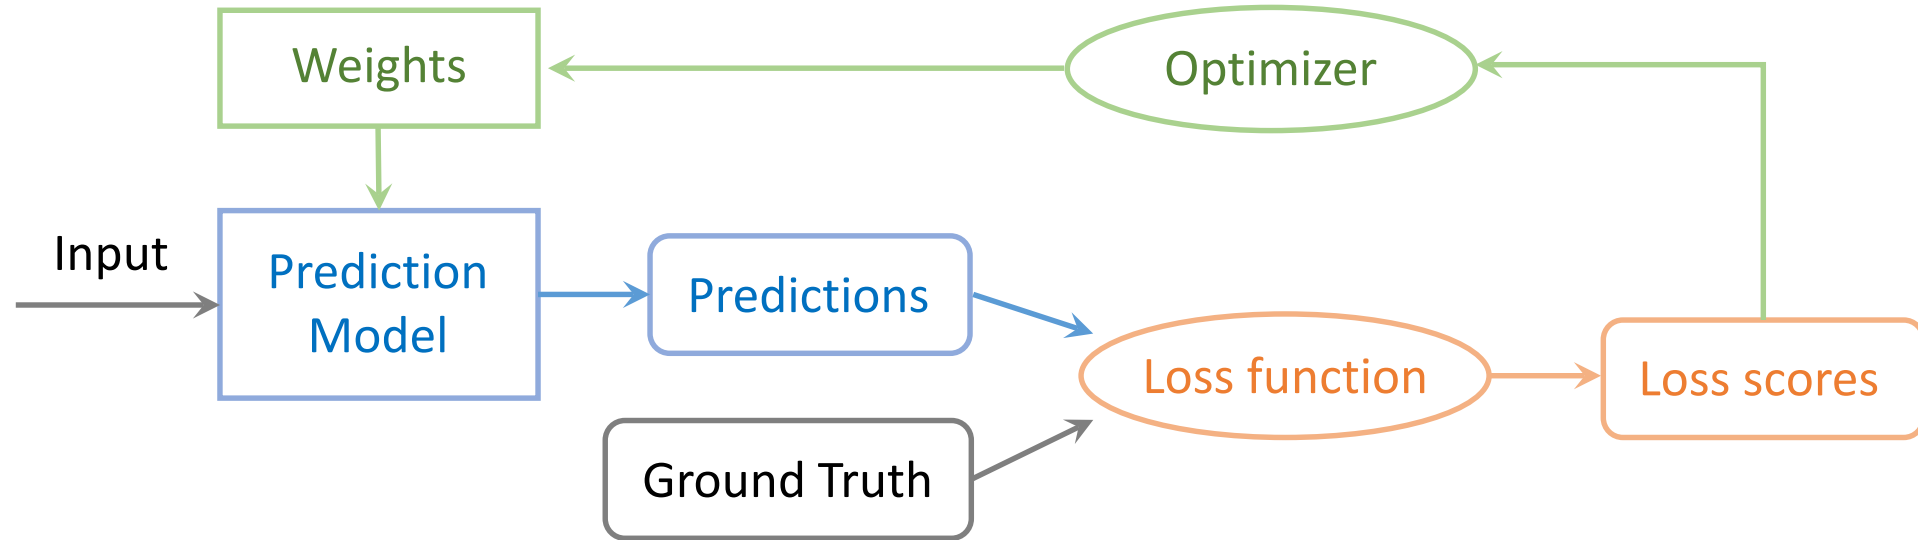




# Outline

- Optimization
  - Gradient Descent
  - Backpropagation
- Neural Networks
  - Input/Hidden/Output Layers
  - Keras for Training Neural Networks

# Training Pipeline



- The training pipeline consists of choosing the **prediction model**, the **loss function**, and the **optimizer**.
- Once these choices are made, we can feed the input data and labels to start the training process.

# Gradient Descent

```
keras.optimizers.SGD(lr=0.001)
```

- Consider the following sequence

$$\mathbf{w}_{n+1} = \mathbf{w}_n - \alpha \nabla f(\mathbf{w}_n), n \geq 0$$


- Since  $-\nabla f(\mathbf{w}_n)$  is the negative gradient at  $\mathbf{w}_n$ , we have

$$f(\mathbf{w}_0) \geq f(\mathbf{w}_1) \geq f(\mathbf{w}_2) \geq \dots$$

- The scalar  $\alpha$  is called the **learning rate**.
- How to properly select the learning rate?
- Try the Adam optimizer

```
keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999)
```

# Gradient Descent (cont.)




[slide credit: Hung-yi Lee]

# Backpropagation

- $f(x, y, z) = (x + y)z$
- Let  $x = -2, y = 5, z = -4$
- Define  $q = x + y, \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$
- Then  $f = qz, \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$


Calculate  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



# Backpropagation (cont.)

- $f(x, y, z) = (x + y)z$
- Let  $x = -2, y = 5, z = -4$
- Define  $q = x + y, \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$
- Then  $f = qz, \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$

Calculate  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$




Chain Rule

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

Upstream gradient      Local gradient

# Backpropagation (cont.)



[slide credit: Stanford CS231n]



# Neural Networks

- Linear Model:  $f(\mathbf{x}) = \mathbf{W}\mathbf{x}$
- 2-layer Neural Network:


$$f(\mathbf{x}) = f_2(\text{ReLU}(f_1(\mathbf{x}))) = \mathbf{W}_2 \max(0, \mathbf{W}_1 \mathbf{x})$$
$$f_1(\mathbf{x}) = \mathbf{W}_1 \mathbf{x}, \quad f_2(\mathbf{h}) = \mathbf{W}_2 \mathbf{h}, \quad \text{ReLU}(\mathbf{h}) = \max(0, \mathbf{h})$$

# Neural Networks (cont.)

- 2-layer Neural Network:

$$f(\mathbf{x}) = f_2(\text{ReLU}(f_1(\mathbf{x}))) = \mathbf{W}_2 \max(0, \mathbf{W}_1 \mathbf{x})$$

$$f_1(\mathbf{x}) = \mathbf{W}_1 \mathbf{x}, \quad f_2(\mathbf{h}) = \mathbf{W}_2 \mathbf{h}, \quad \text{ReLU}(\mathbf{h}) = \max(0, \mathbf{h})$$




- Take CIFAR10 for example
- $\dim(\mathbf{x})$  is  $32 \times 32 \times 3 = 3072$
- $\dim(\mathbf{s})$  is 10 (number of classes)
- $\dim(\mathbf{h})$  is manually defined

# Neural Network (cont.)


- 3-layer Neural Network:

$$f(\mathbf{x}) = f_3(\text{ReLU}(f_2(\text{ReLU}(f_1(\mathbf{x})))) = \mathbf{W}_3 \max(0, \mathbf{W}_2 \max(0, \mathbf{W}_1 \mathbf{x}))$$

$$f_1(\mathbf{x}) = \mathbf{W}_1 \mathbf{x}, \quad f_2(\mathbf{h}_1) = \mathbf{W}_2 \mathbf{h}_1, \quad f_3(\mathbf{h}_2) = \mathbf{W}_3 \mathbf{h}_2, \quad \text{ReLU}(\mathbf{h}) = \max(0, \mathbf{h})$$



# Keras: Creating Models via Sequential()




```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(32, 32, 3)),
    keras.layers.Dense(512, activation=tf.nn.relu),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dense(10, activation=tf.nn.softmax)])
```

Hyperparameters:


- The number of layers
- The number of neurons per
- The activation functions

# Activation Functions

Sigmoid:  $\sigma(x) = \frac{1}{1+e^{-x}}$




ReLU:  $\max(0, x)$



- Sigmoid leads to the vanishing gradient problem (seldom used).
- ReLU is now widely used in the architecture design of neural networks.

# Sigmoid


Sigmoid:  $\sigma(x) = \frac{1}{1+e^{-x}}$




## Problems:

1. The saturated neurons cause the **vanishing gradient** problem.
2. Sigmoid outputs are not zero-centered.
3.  $\exp()$  is a bit computationally expensive

# Vanishing Gradient Problem



Sigmoid:  $\sigma(x) = \frac{1}{1+e^{-x}}$



[slide credit: Stanford CS231n]

# All Positive Inputs

- What happens when the inputs to a layer are all positive?
- What can we say about the gradients with respect to  $\mathbf{w}$ ?
- Suppose the loss is  $L(f(\mathbf{x}))$  with  $f(\mathbf{x}) = \sum_i w_i x_i + b$


$$\frac{\partial L}{\partial \mathbf{w}} = \frac{\partial L}{\partial f} \frac{\partial f}{\partial \mathbf{w}} \quad \frac{\partial f}{\partial \mathbf{w}} = \left[ \frac{\partial f}{\partial w_1} \quad \cdots \quad \frac{\partial f}{\partial w_n} \right]^T = [x_1 \quad \cdots \quad x_n]^T$$

- Since  $x_i > 0$ , the gradient  $\frac{\partial L}{\partial \mathbf{w}}$  always has the same sign as  $\frac{\partial L}{\partial f}$  (all positive or all negative).



# ReLU

ReLU:  $\max(0, x)$





## Pros:

1. Does not saturate
2. Computationally efficient

## Cons:

1. Not zero-entered output

# Relation to Biological Neurons



# Keras: model.summary()

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 3072)	0
<hr/>		
dense (Dense)	(None, 512)	1573376
<hr/>		
dense_1 (Dense)	(None, 128)	65664
<hr/>		
dense_2 (Dense)	(None, 10)	1290
=====		
Total params: 1,640,330		
Trainable params: 1,640,330		
Non-trainable params: 0		

```
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(32, 32, 3)),
    keras.layers.Dense(512, activation=tf.nn.relu),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dense(10, activation=tf.nn.softmax)])
model.summary()
```

# Keras: Training Pipeline

```
model = create_model()  
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])  
model.fit(train_images, train_labels, epochs=30)
```

```
Epoch 1/30 - 7s 140us/sample - loss: 1.8849 - acc: 0.3227  
Epoch 2/30 - 7s 133us/sample - loss: 1.6746 - acc: 0.3999  
Epoch 3/30 - 7s 131us/sample - loss: 1.5990 - acc: 0.4297  
Epoch 4/30 - 7s 133us/sample - loss: 1.5544 - acc: 0.4458  
Epoch 5/30 - 7s 131us/sample - loss: 1.5119 - acc: 0.4622  
Epoch 6/30 - 7s 132us/sample - loss: 1.4852 - acc: 0.4726  
Epoch 7/30 - 7s 132us/sample - loss: 1.4656 - acc: 0.4780  
⋮  
Epoch 28/30 - 7s 130us/sample - loss: 1.2196 - acc: 0.5629  
Epoch 29/30 - 6s 130us/sample - loss: 1.2128 - acc: 0.5637  
Epoch 30/30 - 6s 130us/sample - loss: 1.2103 - acc: 0.5634
```

# Keras Callbacks

- `model.fit(train_images, train_labels, epochs=30, callbacks = [cp_callback, tb_callback])`

## TensorBoard


```
tb_callback = TensorBoard(log_dir='./logs',  
                           histogram_freq=0,  
                           write_graph=True,  
                           write_grads=True,  
                           write_images=True,  
                           embeddings_freq=0,  
                           embeddings_layer_names=None,  
                           embeddings_metadata=None)
```

## Checkpoint


```
checkpoint_path = "training_cifar10/cp-{epoch:04d}.ckpt"  
cp_callback = keras.callbacks.ModelCheckpoint(  
    checkpoint_path,  
    save_weights_only=True,  
    verbose=1,  
    period=10)
```

# Keras: TensorBoard

epoch\_loss



epoch\_acc



# Keras: Checkpoint

```
checkpoint_path = "training_cifar10/cp-{epoch:04d}.ckpt"
loss, acc = model.evaluate(test_images, test_labels)

for epoch in [10, 20, 30]:
    latest = checkpoint_path.format(epoch=epoch)
    model.load_weights(latest)
    loss, acc = model.evaluate(test_images, test_labels)
```

名稱	大小
checkpoint	1 KB
cp-0010.ckpt.data-00000-of-00001	6,412 KB
cp-0010.ckpt.index	1 KB
cp-0020.ckpt.data-00000-of-00001	6,412 KB
cp-0020.ckpt.index	1 KB
cp-0030.ckpt.data-00000-of-00001	6,412 KB
cp-0030.ckpt.index	1 KB

## Random Weights

1s 86us/sample - loss: 2.5282 - acc: 0.0875

## Epoch 10 checkpoint

1s 60us/sample - loss: 1.5094 - acc: 0.4586

## Epoch 20 checkpoint

1s 60us/sample - loss: 1.4365 - acc: 0.4936

## Epoch 30 checkpoint

1s 59us/sample - loss: 1.4265 - acc: 0.5045